



# Basic R Programming

Veerasak Kritsanapraphan - Software Park Thailand

# Data Science Series

- Data Science for Business
- Practical Data Science with Rapid Miner
- Basic R Programming
- Fundamental Data Science for Data Scientist

3

## Agenda

### Day 1

- Basics of R
- Basics of RStudio
- Basic data types and operations: numbers, characters and composites
- Vectors, creating sequences, common functions
- Importing Data
- Manipulating data
- Lists
- Functions
- if-else statements

4



## Day 2

- Data Wrangling
- Loops
- Apply functions
- Data summaries
- Visualization on Standard Graphics

5



## Day 3

- Visualization on ggplot2
- Hypothesis Testing
- Anova
- Linear Regression

6

# About myself?

- Graduated from San Francisco State University in Master of Science in Computer Information System, 1997
- PhD. Candidate at Chulalongkorn University, research focus on Data Science, Big Data, Mobile Computing and Internet of Thing (IOT)
- Chief Technology Innovation Officer at Greenline Synergy (Subsidiary of BDMS - Bangkok Hospital Group)
- Instructor for Software Park in Data Science, Requirement Discovery and Practical Enterprise Integration
- Instructor for Chulalongkorn University in Data Mining

7

## Slide and Sample Data

<https://github.com/vkrit/r-programming>



8

# What is R?

- R is a system for statistical computation and graphics.
- It is heavily influenced by the **S** language
- **R** was initially written by **Ross Ihaka** and Robert Gentleman at the Department of Statistics of the University of Auckland in Auckland, New Zealand.
- The “**R Core Team**” maintain the source code for the software and release regular updates

9

# What is R?

- In addition, the R project is added to by many of its users, who write source code for many different types of analytical procedures
- Everything from analytical chemistry to epidemiology to linguistics
- Currently 6,000+ different user--written libraries available

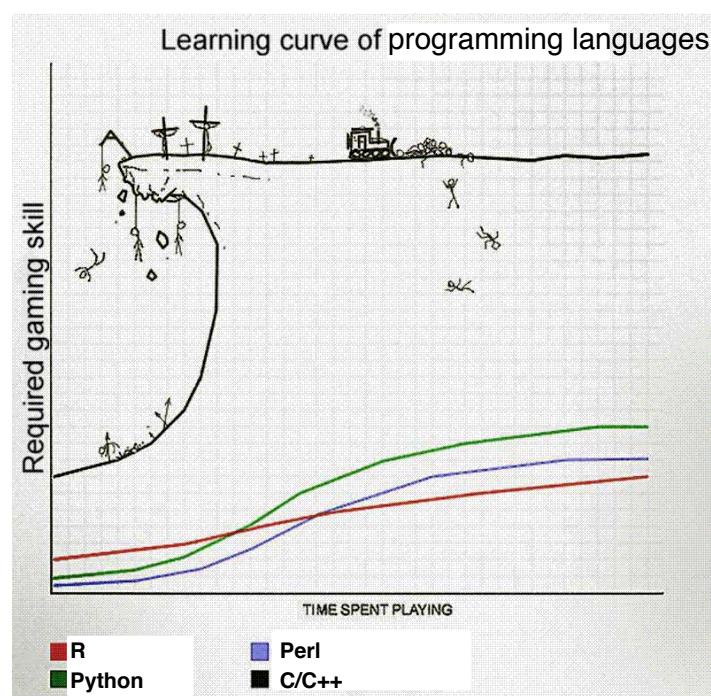
10

# Why use R?

- R is Open-Source Software
- Many built-in functions and installable packages that will cover nearly every possible need
- R is an interpreted language
  - Code doesn't have to be compiled
  - Interactive console makes testing and debugging easy
- Cons to using R
  - Slower than compiled languages
  - Can have runtime errors

11

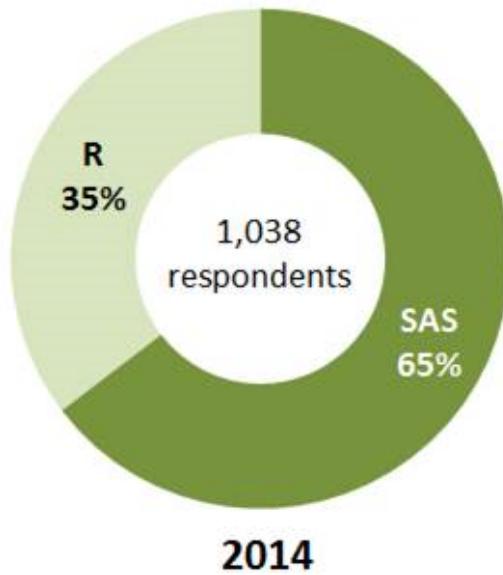
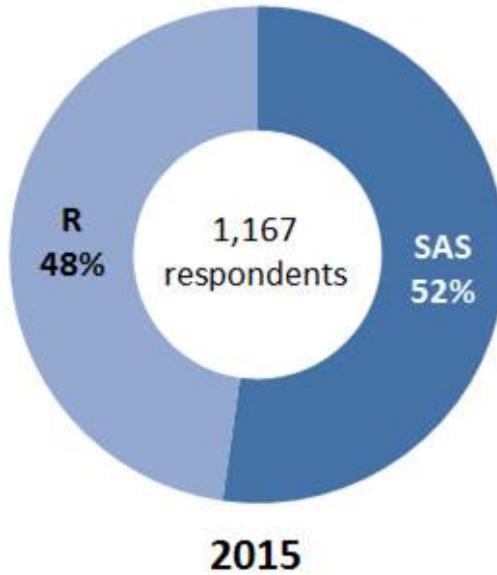
# Why use R?



12

<http://www.vayaputra.es/wordpress/wp-content/uploads/2009/02/2rmqi6o.gif>

# R vs SAS



13

## Tools



**R**  
[www.r-project.org](http://www.r-project.org)

The engine\*



**RStudio**  
[www.rstudio.org](http://www.rstudio.org)

The pretty face\*\*

\* Many alternatives exist. Smallest learning curve.  
\*\* A few alternatives exist. This happens to be the easiest at the moment.

14

# RStudio Features

- Code completion
- Command history search
- Command history to R script / file
- Function extraction from Rscript
- Sweave and Knitr support

15

# Introducing R-Studio

The screenshot shows the RStudio interface with the following components:

- Top Bar:** File, Edit, Code, View, Project, Workspace, Plots, Tools, Help.
- Toolbar:** Includes icons for file operations like Open, Save, Print, and a "Go to file/function" search bar.
- Script Editor:** Shows an R script named "diamondPricing.R" with code for loading ggplot2, source plots/formatPlot.R, and creating a qplot of Price vs Carat by Clarity.
- Console:** Displays the output of the R code, including summary statistics for diamonds and the resulting "Diamond Pricing" plot.
- Workspace:** Shows the dataset "diamonds" has 53940 obs. of 10 variables, and lists objects: avesize, clarity, p, and format.plot.
- Plots:** A scatter plot titled "Diamond Pricing" showing Price (Y-axis, 0 to 15000) versus Carat (X-axis, 0.0 to 3.5). The points are colored by Clarity levels: I1 (red), SI2 (orange), SI1 (yellow), VS2 (green), VS1 (light green), VVS2 (blue), VVS1 (purple), and IF (pink).

16

# RStudio Panes

Rstudio has 4 main windows ('panes'):

- **The Source pane:** create a file that you can save and run later
- **The Console pane:** type or paste in commands to get output from R
- **The Workspace/History pane:** see a list of variables or previous commands
- **The Files/Plots/Packages/Help pane:** see plots, help pages, and other items in this window.

17

## Source and Console Panes

**Source window:** Create a file here, so that you can save and run it later (or turn in as homework)

Any non-command line  
should start with a #

**Console:** Type or paste commands in here to get results from R

If you are loading a data file, you will need to be in the correct directory

> denotes that R is waiting for a command

+ denotes that R is waiting for you to finish the previous command (not shown here)

The screenshot shows the RStudio interface with two main panes: the Source pane and the Console pane. The Source pane contains a script named 'hw1.R' with the following content:

```
1 #  
2 # David Choi  
3 # HW 1  
4 #  
5 # Here are the steps to accomplish problem 1:  
6 x = sqrt(5)  
7 y = sqrt(2)  
8 z = x + y  
9 z
```

The Console pane shows the following R session:

```
11:1 [1] <Top Level>  
Console ~/Dropbox/teaching/DataVis mini 4/homework 1/  
type q to quit.  
[Workspace loaded from ~/.RData]  
> setwd("~/Dropbox/teaching/DataVis mini 4/homework 1")  
> x = sqrt(5)  
> y = sqrt(2)  
> z = x + y  
> z  
[1] 3.650282  
>
```

18

# Console Panes

## Variables

- Save the results of a command in memory by giving it a name:

`z` uses `x` and `y`:

Values are not linked;  
updating `x` doesn't  
change `z`

Redefine `z` using the new  
value for `x`

Use " " or ' ' to distinguish  
between variable names  
and regular text

Note:  
`x <- 3`: same as `x=3`

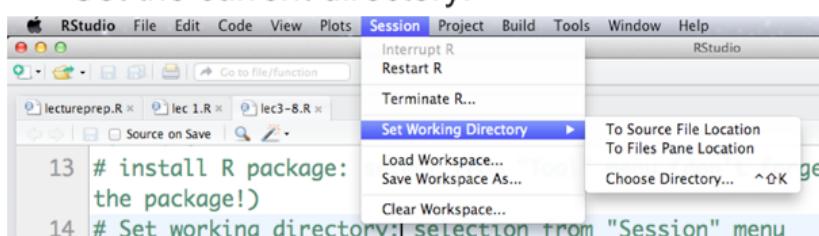
```
Console ~/ ↵
> x = 3
> y = 2
> x           ← Create x and y
[1] 3
> y
[1] 2
> z = sqrt(x^2+y^2)
[1] 3.605551
> x = 5
> z
[1] 3.605551
> z = sqrt(x^2+y^2)
[1] 5.385165
> str = 'Hi there'
> str
[1] "Hi there"
> str2 = "It's cold outside"
> str2
[1] "It's cold outside"
```

19

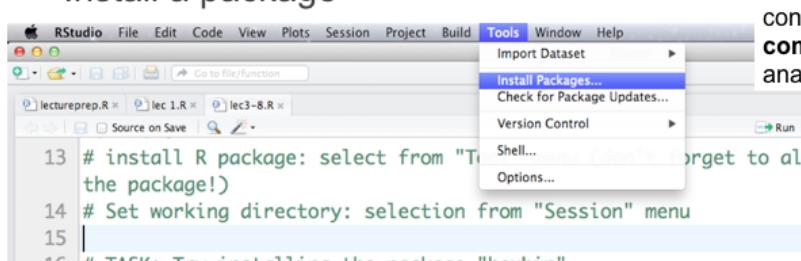
# RStudio Menu

## Two helpful menu items in Rstudio

- Set the current directory:



- Install a package



Packages are extensions to R, containing **new commands** for analysis or graphics

20

# R Markdown

```
example.Rmd x
ABC Knit HTML Chunks
1 Header 1
2 -----
3 This is an R Markdown document. Markdown is a
| simple formatting syntax for authoring web pages.
4
5 Use an asterisk mark, to provide emphasis such as
| *italics* and **bold**.
6
7 Create lists with a dash:
8 - Item 1
9 - Item 2
10 - Item 3
11
12 You can write `in-line` code with a back-tick.
13
14 ...
15 Code blocks display
16 with fixed-width font
17 ...
18
19 > Blockquotes are offset
20
```

Header 1

This is an R Markdown document. Markdown is a simple formatting syntax for authoring web pages.

Use an asterisk mark, to provide emphasis such as *italics* and **bold**.

Create lists with a dash:

- Item 1
- Item 2
- Item 3

You can write `in-line` code with a back-tick.

Code blocks display  
with fixed-width font

Blockquotes are offset

21

# R Markdown

```
chunks.Rmd x
ABC Knit HTML Chunks
1 R Code Chunks
2 -----
3
4 With R Markdown, you can insert R code
| chunks including plots:
5
6 ````{r qplot, fig.width=4, fig.height=3,
| message=FALSE}
7 # quick summary and plot
8 library(ggplot2)
9 summary(cars)
10 qplot(speed, dist, data=cars) +
11     geom_smooth()
12 ...
13 |
```

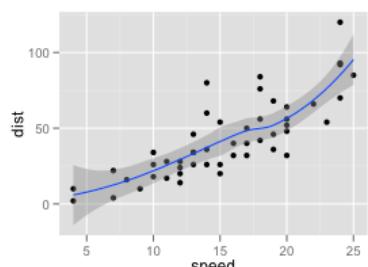
R Code Chunks

With R Markdown, you can insert R code chunks including plots:

```
# quick summary and plot
library(ggplot2)
summary(cars)
```

```
##      speed         dist
## Min.   : 4.0   Min.   : 2
## 1st Qu.:12.0   1st Qu.: 26
## Median :15.0   Median : 36
## Mean   :15.4   Mean   : 43
## 3rd Qu.:19.0   3rd Qu.: 56
## Max.   :25.0   Max.   :120
```

```
qplot(speed, dist, data = cars) + geom_smooth()
```



22

# Workshop 1: Hello World

Open RStudio on your machine

- File > New File > R Markdown ...
- Change summary(cars) in the first code block to print("Hello world!")
- Click Knit HTML to produce an HTML file.
- Save your Rmd file as helloworld.Rmd

23

## Data Types

You'll encounter different kinds of data types

- **Booleans** Direct binary values: TRUE or FALSE in R
- **Integers**: whole numbers (positive, negative or zero)
- **Characters** fixed-length blocks of bits, with special coding; strings = sequences of characters
- **Floating point numbers**: a fraction (with a finite number of bits) times an exponent, like  $1.87 * 10^6$
- **Missing** or ill-defined values: NA, NaN, etc.

24

# Operators (Functions)

Command	Description
<code>+,-,*,\</code>	add, subtract, multiply, divide
<code>^</code>	raise to the power of
<code>%%</code>	remainder after division (ex: <code>8 %% 3 = 2</code> )
<code>( )</code>	change the order of operations
<code>log(), exp()</code>	logarithms and exponents (ex: <code>log(10) = 2.302</code> )
<code>sqrt()</code>	square root
<code>round()</code>	round to the nearest whole number (ex: <code>round(2.3) = 2</code> )
<code>floor(), ceiling()</code>	round down or round up
<code>abs()</code>	absolute value

25

## Operators

```
7 + 5 # Addition
```

```
[1] 12
```

```
7 - 5 # Subtraction
```

```
[1] 2
```

```
7 * 5 # Multiplication
```

```
[1] 35
```

```
7 ^ 5 # Exponentiation
```

```
[1] 16807
```

26



```
7 / 5 # Division
```

```
[1] 1.4
```

```
7 %% 5 # Modulus
```

```
[1] 2
```

```
7 %/% 5 # Integer division
```

```
[1] 1
```

27

# Operators

**Comparisons** are also **binary operators**; they take two objects, like numbers. and give a Boolean

```
7 > 5
```

```
[1] TRUE
```

```
7 < 5
```

```
[1] FALSE
```

```
7 >= 7
```

```
[1] TRUE
```

```
7 <= 5
```

```
[1] FALSE
```

28



```
7 == 5
```

```
[1] FALSE
```

```
7 != 5
```

```
[1] TRUE
```

29

## Boolean Operators

Basically “and” and “or”:

```
(5 > 7) & (6*7 == 42)
```

```
[1] FALSE
```

```
(5 > 7) | (6*7 == 42)
```

```
[1] TRUE
```

30

# Basic Math / Basic Logic

- + addition
- subtraction
- \* multiplication
- / division
- % modulus (remainder)
- $\wedge$  to the power

?Arithmetic

- ! NOT
- & bitwise AND
- | bitwise OR
- $\&\&$  short circuit AND
- $\|$  short circuit OR
- $\equiv$  equality
- $\neq$  NOT equality

?Logic

Also try ?Syntax, ?Comparison

31

## More types

`typeof()` function returns the type

`is.foo()` functions return Booleans for whether the argument is of type foo

`as.foo()` (tries to) “cast” its argument to type foo — to translate it sensibly into a foo-type value

32



```
typeof(7)
```

```
[1] "double"
```

```
is.numeric(7)
```

```
[1] TRUE
```

```
is.na(7)
```

```
[1] FALSE
```

33



```
is.character(7)
```

```
[1] FALSE
```

```
is.character("7")
```

```
[1] TRUE
```

```
is.character("seven")
```

```
[1] TRUE
```

```
is.na("seven")
```

```
[1] FALSE
```

34

# Variables

We can give names to data objects; these give us **variables**

A few variables are built in:

```
pi
```

```
[1] 3.141593
```

Variables can be arguments to functions or operators, just like constants:

```
pi*10
```

```
[1] 31.41593
```

```
cos(pi)
```

```
[1] -1
```

35

# Assignment Operator

Most variables are created with the **assignment operator**, `<-` or `=`

```
approx.pi <- 22 / 7  
approx.pi
```

```
[1] 3.142857
```

```
diameter.in.cubits = 10  
approx.pi*diameter.in.cubits
```

```
[1] 31.42857
```

36

# Variables

- Using names and variables makes code: easier to design, easier to debug, less prone to bugs, easier to improve, and easier for others to read
- Avoid “magic constants”; use named variables
- Use descriptive variable names
  - Good: num.students <- 35
  - Bad: ns <- 35

37

# Workspace

What names have you defined values for?

```
ls()
```

```
[1] "approx.pi"           "circumference.in.cubits"  
[3] "diameter.in.cubits"
```

Getting rid of variables:

```
rm("circumference.in.cubits")  
ls()
```

```
[1] "approx.pi"           "diameter.in.cubits"
```

38

# Data Structure : Vector

- Group related data values into one object, a **data structure**
- A **vector** is a sequence of values, all of the same type
- `c()` function returns a vector containing all its arguments in order

```
students <- c("Sean", "Louisa", "Frank", "Farhad", "Li")
midterm <- c(80, 90, 93, 82, 95)
```

- Typing the variable name at the prompt causes it to display

```
students
[1] "Sean"    "Louisa"   "Frank"    "Farhad"   "Li"
```

39

# Indexing Vector

- `vec[1]` is the first element, `vec[4]` is the 4th element of `vec`

```
students
[1] "Sean"    "Louisa"   "Frank"    "Farhad"   "Li"
students[4]
[1] "Farhad"
```

- `vec[-4]` is a vector containing all but the fourth element

```
students[-4]
[1] "Sean"    "Louisa"   "Frank"    "Li"
```

40

# Vector Arithmetic

Operators apply to vectors “pairwise” or “elementwise”:

```
final <- c(78, 84, 95, 82, 91) # Final exam scores  
midterm # Midterm exam scores  
  
[1] 80 90 93 82 95  
  
midterm + final # Sum of midterm and final scores  
  
[1] 158 174 188 164 186  
  
(midterm + final)/2 # Average exam score  
  
[1] 79 87 94 82 93  
  
course.grades <- 0.4*midterm + 0.6*final # Final course grade  
course.grades  
  
[1] 78.8 86.4 94.2 82.0 92.6
```

41

# Pairwise Comparison

Is the final score higher than the midterm score?

```
midterm  
  
[1] 80 90 93 82 95  
  
final  
  
[1] 78 84 95 82 91  
  
final > midterm  
  
[1] FALSE FALSE TRUE FALSE FALSE
```

Boolean operators can be applied elementwise:

```
(final < midterm) & (midterm > 80)  
  
[1] FALSE TRUE FALSE FALSE TRUE
```

42

# Functions on Vectors

Command	Description
sum(vec)	sums up all the elements of vec
mean(vec)	mean of vec
median(vec)	median of vec
min(vec), max(vec)	the largest or smallest element of vec
sd(vec), var(vec)	the standard deviation and variance of vec
length(vec)	the number of elements in vec
pmax(vec1, vec2), pmin(vec1, vec2)	example: pmax(quiz1, quiz2) returns the higher of quiz 1 and quiz 2 for each student
sort(vec)	returns the vec in sorted order
order(vec)	returns the index that sorts the vector vec
unique(vec)	lists the unique elements of vec
summary(vec)	gives a five-number summary
any(vec), all(vec)	useful on Boolean vectors

43

# Function on Vectors

```
course.grades
```

```
[1] 78.8 86.4 94.2 82.0 92.6
```

```
mean(course.grades) # mean grade
```

```
[1] 86.8
```

```
median(course.grades)
```

```
[1] 86.4
```

```
sd(course.grades) # grade standard deviation
```

```
[1] 6.625708
```

44

# More on Functions

```
sort(course.grades)
```

```
[1] 78.8 82.0 86.4 92.6 94.2
```

```
max(course.grades) # highest course grade
```

```
[1] 94.2
```

```
min(course.grades) # lowest course grade
```

```
[1] 78.8
```

45

# Referencing elements of Vectors

```
students
```

```
[1] "Sean"    "Louisa"  "Frank"   "Farhad" "Li"
```

Vector of indices:

```
students[c(2,4)]
```

```
[1] "Louisa" "Farhad"
```

Vector of negative indices

```
students[c(-1,-3)]
```

```
[1] "Louisa" "Farhad" "Li"
```

46

# More on Referencing

`which()` returns the TRUE indexes of a Boolean vector:

```
course.grades
```

```
[1] 78.8 86.4 94.2 82.0 92.6
```

```
a.threshold <- 90 # A grade = 90% or higher  
course.grades >= a.threshold # vector of booleans
```

```
[1] FALSE FALSE TRUE FALSE TRUE
```

```
a.students <- which(course.grades >= a.threshold) # Applying which()  
a.students
```

```
[1] 3 5
```

```
students[a.students] # Names of A students
```

```
[1] "Frank" "Li"
```

47

# Workshop 2

Using the same Rmarkdown file from Workshop 1, do the following:-

1) Creating sequences

**: Colon operator:**

```
1:10 # Numbers 1 to 10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
127:132 # Numbers 127 to 132
```

```
## [1] 127 128 129 130 131 132
```

48

# Workshop 2

**seq** function: `seq(from, to, by)`

```
seq(1,10,1) # Numbers 1 to 10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
seq(1,10,2) # Odd numbers from 1 to 10
```

```
## [1] 1 3 5 7 9
```

```
seq(2,10,2) # Even numbers from 2 to 10
```

```
## [1] 2 4 6 8 10
```

49

## Workshop 2 : Use RStudio to answer questions below

- (a) Use : to output the sequence of numbers from 3 to 12
- (b) Use seq() to output the sequence of numbers from 3 to 30 in increments of 3
- (c) Save the sequence from (a) as a variable x, and the sequence from (b) as a variable y. Output their product x\*y

50

# Difference between = and <-

- The operators <- and = assign into the environment in which they are evaluated.
- The operator <- can be used anywhere, whereas the operator = is only allowed at the top level (e.g., in the complete expression typed at the command prompt) or as one of the subexpressions in a braced list of expressions.

```
matrix(1,nrow=2)
```

```
matrix(1,nrow<-2)
```

51

# Reading and Getting Data into R

## Combine Command

```
c(1, 2, 3, 4)  
c(item1, item2, item3, item4)  
c("item1", "item2", "item3")
```

## Scan Command

```
our.data = scan()  
  
scan(what = 'character')  
data5 = scan(sep = ',', what = 'char')  
data6 = scan(file = 'data.txt')
```

## Working Directory

```
getwd()
```

52

# Reading Bigger Data Files

```
read.csv()  
read.csv(file, sep = ',', header = TRUE, row.names)  
fw = read.csv(file.choose())
```

```
my.ssv = read.table(file.choose(), header = TRUE)  
my.tsv = read.delim(file.choose())  
my.tsv = read.csv(file.choose(), sep = '\t')  
my.tsv = read.table(file.choose(), header = TRUE,  
sep = '\t')
```

53

# Importing Data

- To import tabular data into R, we use the `read.table()` command

```
1 survey <- read.table("survey_data.csv", header=TRUE, sep=",")  
2  
3
```

- Let's parse this command one component at a time
  - The data is in a file called `survey_data.csv`, which is an online file
  - The file contains a header as its first row
  - The csv format means that the data is comma-separated, so `sep=" , "`
- Could've also used `read.csv()`, which is just `read.table()` with the preset `sep=" , "`

54

# Exploring the Data

- R imports data into a `data.frame` object

```
class(survey)
```

```
[1] "data.frame"
```

- To view the first few rows of the data, use `head()`

```
head(survey, 3)
```

	Program	PriorExp	Rexperience	OperatingSystem	TVhours
1	MISM	Some experience	Never used	Windows	1
2	Other	Some experience	Basic competence	Windows	8
3	MISM	Extensive experience	Basic competence	Windows	4

Editor

1 Microsoft Word  
2 Microsoft Word  
3 Microsoft Word

- `head(data.frame, n)` returns the first n rows of the data frame
- In the Console, you can also use `View(survey)` to get a spreadsheet view

55

# Simple Summary

- Use the `str()` function to get a simple summary of your data set

```
str(survey)
```

```
'data.frame': 31 obs. of 6 variables:  
 $ Program      : Factor w/ 3 levels "MISM", "Other", ... : 1 2 1 3 2 3 2 1 3 3 ...  
 $ PriorExp     : Factor w/ 3 levels "Extensive experience", ... : 3 3 1 2 2 3 3 3 3 ...  
 $ Rexperience   : Factor w/ 3 levels "Basic competence", ... : 3 1 1 2 3 3 3 3 3 3 ...  
 ...  
 $ OperatingSystem: Factor w/ 2 levels "Mac OS X", "Windows": 2 2 2 2 2 1 1 2 1 2 ...  
 ...  
 $ TVhours       : int 1 8 4 5 3 10 0 10 15 4 ...  
 $ Editor        : Factor w/ 1 level "Microsoft Word": 1 1 1 1 1 1 1 1 1 1 ...
```

- This says that `TVhours` is a numeric variable, while all the rest are factors (categorical)

56

# Another Summary

```
summary(survey)
```

Program	PriorExp	Rexperience
MISM :10	Extensive experience : 5	Basic competence : 2
Other: 6	Never programmed before:12	Installed on machine: 7
PPM :15	Some experience :14	Never used :22

OperatingSystem	TVhours	Editor
Mac OS X:14	Min. : 0.000	Microsoft Word:31
Windows :17	1st Qu.: 1.000	
	Median : 4.000	
	Mean : 5.742	
	3rd Qu.: 9.000	
	Max. :40.000	

57

# Data Frames

- 2 Dimensional Objects, it has rows and columns. R treats the columns as separate samples or variables, rows represent the replicates or observations.
- To see what an R object is made up of, you can use `attributes()`

```
attributes(survey)
```

```
$names
[1] "Program"          "PriorExp"        "Rexperience"      "OperatingSystem"
[5] "TVhours"          "Editor"

$class
[1] "data.frame"

$row.names
 [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
[24] 24 25 26 27 28 29 30 31
```

An R **data frame** is a *list* whose columns you can refer to by *name* or *index*

58

# Data Frame dimensions

- We can use `nrow()` and `ncol` to determine the number of survey responses and the number of survey questions

```
nrow(survey) # Number of rows (responses)  
[1] 31  
  
ncol(survey) # Number of columns (questions)  
[1] 6
```

- When writing reports, you will often want to say how large your sample size was
- To do this *inline*, use the syntax:

```
`r nrow(survey)`
```

- This allows us to write “31 students responded to the survey”, and have the number displayed automatically change when `nrow(survey)` changes.

59

# Indexing Data Frame

- There are many different ways of indexing the same piece of a data frame

```
survey[["Program"]] # "Program" element  
[1] MISM Other MISM PPM Other PPM Other MISM PPM PPM MISM  
[12] PPM PPM PPM PPM PPM MISM MISM MISM PPM PPM  
[23] MISM Other PPM PPM MISM PPM Other Other MISM  
Levels: MISM Other PPM  
  
survey$Program # "Program" element  
[1] MISM Other MISM PPM Other PPM Other MISM PPM PPM MISM  
[12] PPM PPM PPM PPM PPM MISM MISM MISM PPM PPM  
[23] MISM Other PPM PPM MISM PPM Other Other MISM  
Levels: MISM Other PPM  
  
survey[,1] # Data from 1st column  
[1] MISM Other MISM PPM Other PPM Other MISM PPM PPM MISM  
[12] PPM PPM PPM PPM PPM MISM MISM MISM PPM PPM  
[23] MISM Other PPM PPM MISM PPM Other Other MISM  
Levels: MISM Other PPM
```

60

# More Indexing

- Note that single brackets and double brackets have different effects

```
survey[["Program"]]
```

```
[1] MISM Other MISM PPM Other PPM Other MISM PPM PPM MISM  
[12] PPM PPM PPM PPM PPM MISM MISM MISM PPM PPM  
[23] MISM Other PPM PPM MISM PPM Other Other MISM  
Levels: MISM Other PPM
```

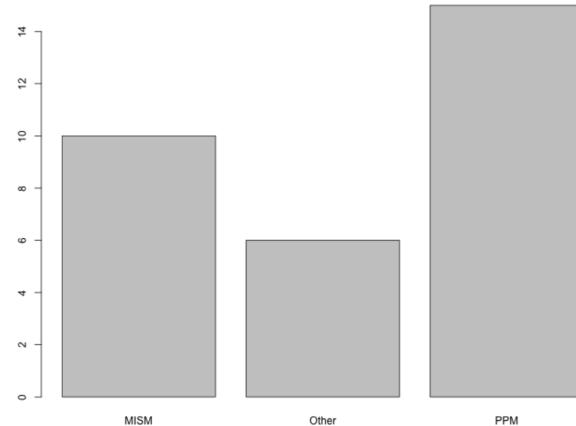
```
survey["Program"] # sub-data frame containing only "Program"
```

```
Program  
1     MISM  
2     Other  
3     MISM  
4     PPM  
5     Other  
6     PPM  
7     Other  
8     MISM  
9     PPM  
10    PPM  
11    MISM  
12    PPM  
13    PPM  
14    PPM
```

61

# Bar Plot (Categorical Data)

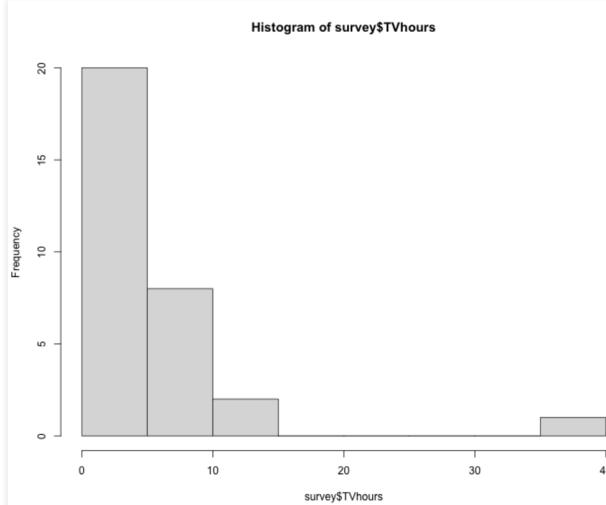
```
plot(survey[["Program"]])
```



62

# Histogram (Continuous Data)

```
hist(survey$TVhours, col="lightgray")
```



63

# Indexing Multiple Columns

```
head(survey[,c(1,5)]) # Data from 1st and 5th columns
```

```
Program TVhours
1    MISM      1
2    Other      8
3    MISM      4
4    PPM       5
5    Other      3
6    PPM      10
```

```
head(survey[c("Program", "Editor")]) # Data from "Program" and "Editor"
```

```
Program          Editor
1    MISM Microsoft Word
2    Other Microsoft Word
3    MISM Microsoft Word
4    PPM  Microsoft Word
5    Other Microsoft Word
6    PPM  Microsoft Word
```

64

# Indexing Row and Column

- Data frames have two dimensions to index across

```
survey[6,] # 6th row
```

```
Program      PriorExp Rexperience OperatingSystem TVhours
6      PPM Some experience    Never used        Mac OS X      10
      Editor
6 Microsoft Word
```

```
survey[6,5] # row 6, column 5
```

```
[1] 10
```

```
survey[6, "Program"] # Program of 6th survey respondent
```

```
[1] PPM
Levels: MISM Other PPM
```

```
survey[["Program"]][6] # Program of 6th survey respondent
```

```
[1] PPM
Levels: MISM Other PPM
```

65

# More Indexing

- We can use this operator for indexing

```
survey[1:3,] # equivalent to head(survey, 3)
```

```
Program      PriorExp      Rexperience OperatingSystem TVhours
1      MISM      Some experience    Never used        Windows      1
2      Other      Some experience Basic competence        Windows      8
3      MISM Extensive experience Basic competence        Windows      4
      Editor
1 Microsoft Word
2 Microsoft Word
3 Microsoft Word
```

```
survey[3:5, c(1,5)]
```

```
Program TVhours
3      MISM      4
4      PPM       5
5      Other      3
```

66

# Subsets of Data

- We are often interested in learning something about a specific subset of the data

```
survey[survey$Program=="MISM",] # Data from the MISM students
```

	Program	PriorExp	Rexperience	OperatingSystem
1	MISM	Some experience	Never used	Windows
3	MISM	Extensive experience	Basic competence	Windows
8	MISM	Some experience	Never used	Windows
11	MISM	Extensive experience	Installed on machine	Mac OS X
18	MISM	Extensive experience	Never used	Mac OS X
19	MISM	Extensive experience	Never used	Windows
20	MISM	Some experience	Never used	Windows
23	MISM	Some experience	Never used	Windows
27	MISM	Some experience	Never used	Mac OS X
31	MISM	Some experience	Installed on machine	Windows
	TVhours	Editor		
1	1	Microsoft Word		
3	4	Microsoft Word		
8	10	Microsoft Word		
11	0	Microsoft Word		
18	3	Microsoft Word		
19	0	Microsoft Word		
20	1	Microsoft Word		
23	3	Microsoft Word		
27	0	Microsoft Word		
31	0	Microsoft Word		

67

## More subsets example

- Let's pull all of the PPM students who have never used R before

```
survey[survey$Program=="PPM" & survey$Rexperience=="Never used",]
```

	Program	PriorExp	Rexperience	OperatingSystem	TVhours
6	PPM	Some experience	Never used	Mac OS X	10
9	PPM	Never programmed before	Never used	Mac OS X	15
10	PPM	Extensive experience	Never used	Windows	4
12	PPM	Never programmed before	Never used	Windows	0
13	PPM	Some experience	Never used	Mac OS X	10
14	PPM	Never programmed before	Never used	Mac OS X	4
15	PPM	Some experience	Never used	Windows	10
16	PPM	Some experience	Never used	Mac OS X	2
22	PPM	Never programmed before	Never used	Windows	7
25	PPM	Never programmed before	Never used	Mac OS X	6
		Editor			
6	Microsoft Word				
9	Microsoft Word				
10	Microsoft Word				
12	Microsoft Word				
13	Microsoft Word				
14	Microsoft Word				
15	Microsoft Word				
16	Microsoft Word				
22	Microsoft Word				
25	Microsoft Word				

68

# Subset

- When the subset conditions get long or messy, it is preferable to use the `subset()` function
- Here's an example of selecting the OperatingSystem and TVhours responses from all of the students who are either in PPM or Other and who listed their R experience as "Basic competence".

```
subset(survey, select=c("OperatingSystem", "TVhours"), subset=(Program == "PPM"  
| Program == "Other") & Rexperience == "Basic competence")
```

```
OperatingSystem TVhours  
2 Windows 8
```

69

# Calculation from subset

```
mean(survey$TVhours[survey$Program == "PPM"]) # Average time PPM's spent watching  
TV
```

```
[1] 8.8
```

```
mean(survey$TVhours[survey$Program == "MISM"]) # Average time MISM's spent  
watching TV
```

```
[1] 2.2
```

```
mean(survey$TVhours[survey$Program == "Other"]) # Average time "Others" spent  
watching TV
```

```
[1] 4
```

70

# Coding Style

- Coding style (and code commenting) will become increasingly more important as we get into more advanced and involved programming tasks
- A few R “style guides” exist:
  - <http://r-pkgs.had.co.nz/style.html>
  - Google’s R Style Guide (<https://google.github.io/styleguide/Rguide.xml>)

71

## Example Style Guide

### Assignment operator. USE `<-`

```
student.names <- c("Eric", "Hao", "Jennifer") # Good
student.names = c("Eric", "Hao", "Jennifer") # Bad
```

- Note: When specifying function arguments, only = is valid

```
sort(tv.hours, decreasing=TRUE) # Good
sort(tv.hours, decreasing<-TRUE) # Bad!!
```

72

# Style Guide : Variable Name

- To make code easy to read, debug, and maintain, you should use **concise** but **descriptive** variable names
- Terms in variable names should be separated by `_` or `.`

```
# Accepted
day_one    day.one   day_1    day.1    day1

# Bad
d1      DayOne   dayone

# Can be made more concise:
first.day.of.the.month
```

- Avoid using variable names that are already pre-defined in R

```
# EXTREMELY bad:
c      T      pi     sum    mean
```

73

## Workshop 3: Importing and Indexing data

### Data practice

```
survey <- read.table("survey.csv");
```

Use command to answer these questions?

(a) How many survey respondents are from MISM or Other?

(b) What % of survey respondents are from PPM?

### Index practice

(a) Use `$` notation to pull the `OperatingSystem` column from the survey data

(b) Do the same thing with `[,]` notation, referring to `OperatingSystem` by name

(c) Repeat part (b), this time referring to `OperatingSystem` by column number

74

# More on Data Frame

```
library(MASS)
head(Cars93, 3)
```

	Manufacturer	Model	Type	Min.Price	Price	Max.Price	MPG.city
1	Acura	Integra	Small	12.9	15.9	18.8	25
2	Acura	Legend	Midsize	29.2	33.9	38.7	18
3	Audi	90	Compact	25.9	29.1	32.3	20
	MPG.highway		AirBags	DriveTrain	Cylinders	EngineSize	
1	31		None	Front	4	1.8	
2	25	Driver & Passenger		Front	6	3.2	
3	26	Driver only		Front	6	2.8	
	Horsepower	RPM	Rev.per.mile	Man.trans.avail	Fuel.tank.capacity		
1	140	6300	2890		Yes	13.2	
2	200	5500	2335		Yes	18.0	
3	172	5500	2280		Yes	16.9	
	Passengers	Length	Wheelbase	Width	Turn.circle	Rear.seat.room	
1	5	177	102	68	37	26.5	
2	5	195	115	71	38	30.0	
3	5	180	102	67	37	28.0	
	Luggage.room	Weight	Origin		Make		
1	11	2705	non-USA	Acura	Integra		
2	15	3560	non-USA	Acura	Legend		
3	14	3375	non-USA		Audi 90		

75

## Add column using “Transform”

- `transform()` returns a new data frame with columns modified or added as specified by the function call

```
Cars93.metric <- transform(Cars93,
                           KMPL.city = 0.425 * MPG.city,
                           KMPL.highway = 0.425 * MPG.highway)
tail(names(Cars93.metric))
```

```
[1] "Luggage.room"   "Weight"        "Origin"        "Make"
[5] "KMPL.city"       "KMPL.highway"
```

- Our data frame has two new columns, giving the fuel consumption in km/l

76

# Another Approach

```
KMPL.city.2 <- 0.425 * Cars93$MPG.city  
# Add a new column called KMPL.city.2  
Cars93.metric$KMPL.city.2 <- KMPL.city.2  
tail(names(Cars93.metric))
```

```
[1] "Weight"      "Origin"       "Make"        "KMPL.city"  
[5] "KMPL.highway" "KMPL.city.2"
```

- Let's check that both approaches did the same thing

```
identical(Cars93.metric$KMPL.city, Cars93.metric$KMPL.city.2)
```

```
[1] TRUE
```

77

# Changing level of a factor

```
manufacturer <- Cars93$Manufacturer  
head(manufacturer, 10)
```

```
[1] Acura      Acura      Audi      Audi      BMW       Buick     Buick  
[8] Buick      Buick      Cadillac  
32 Levels: Acura Audi BMW Buick Cadillac Chevrolet Chrylser ... Volvo
```

We'll use the `mapvalues(x, from, to)` function from the `plyr` library.

```
library(plyr)  
  
# Map Chevrolet, Pontiac and Buick to GM  
manufacturer.combined <- mapvalues(manufacturer,  
                                    from = c("Chevrolet", "Pontiac", "Buick"),  
                                    to = rep("GM", 3))
```

```
head(manufacturer.combined, 10)
```

```
[1] Acura      Acura      Audi      Audi      BMW       GM        GM  
[8] GM        GM        Cadillac  
30 Levels: Acura Audi BMW GM Cadillac Chrylser Chrysler Dodge ... Volvo
```

78

# Another example

- A lot of data comes with integer encodings of levels
- You may want to convert the integers to more meaningful values for the purpose of your analysis
- Let's pretend that in the class survey 'Program' was coded as an integer with 1 = MISM, 2 = Other, 3 = PPM
- Here's how we would get back the program codings using the `transform()`, `as.factor()` and `mapvalues()` functions

79

```
survey <- transform(survey, Program = as.factor(mapvalues(Program, c(1, 2, 3),
c("MISM", "Other", "PPM"))))
```

	Program	PriorExp	Rexperience	OperatingSystem
1	MISM	Some experience	Never used	Windows
2	Other	Some experience	Basic competence	Windows
3	MISM	Extensive experience	Basic competence	Windows
4	PPM	Never programmed before	Installed on machine	Windows
5	Other	Never programmed before	Never used	Windows
6	PPM	Some experience	Never used	Mac OS X
	TVhours	Editor		
1	1	Microsoft Word		
2	8	Microsoft Word		
3	4	Microsoft Word		
4	5	Microsoft Word		
5	3	Microsoft Word		
6	10	Microsoft Word		

80

# table() function

- The `table()` function builds **contingency tables** showing counts at each combination of factor levels

```
table(Cars93$AirBags)
```

Driver & Passenger	16	Driver only	43
		None	34

```
table(Cars93$Origin)
```

USA	non-USA
48	45

```
table(Cars93$AirBags, Cars93$Origin)
```

	USA	non-USA
Driver & Passenger	9	7
Driver only	23	20
None	16	18

- Looks like US and non-US cars had about the same distribution of AirBag types

81

# more table()

- When `table()` is supplied a data frame, it produces contingency tables for all combinations of factors

```
head(Cars93[c("AirBags", "Origin")], 3)
```

	AirBags	Origin
1	None	non-USA
2	Driver & Passenger	non-USA
3	Driver only	non-USA

```
table(Cars93[c("AirBags", "Origin")])
```

AirBags		Origin	
		USA	non-USA
Driver & Passenger	9	7	
Driver only	23	20	
None	16	18	

82

# Basic of Lists

A list is a **data structure** that can be used to store **different kinds** of data

- Recall: a vector is a data structure for storing *similar kinds of data*
- To better understand the difference, consider the following example.

```
my.vector.1 <- c("Michael", 165, TRUE) # (name, weight, is.male)
my.vector.1
[1] "Michael" "165"      "TRUE"
typeof(my.vector.1) # All the elements are now character strings!
[1] "character"
```

83

# List vs Vector

```
my.vector.2 <- c(FALSE, TRUE, 27) # (is.male, is.citizen, age)
typeof(my.vector.2)
[1] "double"
```

- Vectors expect elements to be all of the same type (e.g., Boolean, numeric, character)
- When data of different types are put into a vector, the R converts everything to a common type

84

# Lists

- To store data of different types in the same object, we use lists
- Simple way to build lists: use `list()` function

```
my.list <- list("Michael", 165, TRUE)  
my.list
```

```
[[1]]  
[1] "Michael"  
  
[[2]]  
[1] 165  
  
[[3]]  
[1] TRUE
```

```
sapply(my.list, typeof)
```

```
[1] "character" "double"    "logical"
```

85

# Named Elements

```
patient.1 <- list(name="Michael", weight=165, is.male=TRUE)  
patient.1
```

```
$name  
[1] "Michael"  
  
$weight  
[1] 165  
  
$is.male  
[1] TRUE
```

86

# Referencing element inside list

```
patient.1$name # Get "name" element (returns a string)  
[1] "Michael"  
  
patient.1[["name"]] # Get "name" element (returns a string)  
[1] "Michael"  
  
patient.1["name"] # Get "name" slice (returns a sub-list)  
  
$name  
[1] "Michael"  
  
c(typeof(patient.1$name), typeof(patient.1["name"]))  
[1] "character" "list"
```

87

# Function

- We have used a lot of built-in functions: `mean()`, `subset()`, `plot()`, `read.table()`...
- An important part of programming and data analysis is to write custom functions
- Functions help make code **modular**
- Functions make debugging easier
- Remember: this entire class is about applying *functions* to *data*

88

# what is a function?

A function is a machine that turns **input objects** (arguments) into an **output object** (return value) according to a definite rule.

- Let's look at a really simple function

```
addOne <- function(x) {  
  x + 1  
}
```

- x is the **argument** or **input**
- The function **output** is the input x incremented by 1

```
addOne(12)
```

```
[1] 13
```

89

## Another example

- Here's a function that returns a % given a numerator, denominator, and desired number of decimal values

```
calculatePercentage <- function(x, y, d) {  
  decimal <- x / y # Calculate decimal value  
  round(100 * decimal, d) # Convert to % and round to d digits  
}  
  
calculatePercentage(27, 80, 1)
```

```
[1] 33.8
```

- If you're calculating several %'s for your report, you should use this kind of function instead of repeatedly copying and pasting code

90

# Function return a list

- Here's a function that takes a person's full name (FirstName LastName), weight in lb and height in inches and converts it into a list with the person's first name, person's last name, weight in kg, height in m, and BMI.

```
createPatientRecord <- function(full.name, weight, height) {  
  name.list <- strsplit(full.name, split=" ")[[1]]  
  first.name <- name.list[1]  
  last.name <- name.list[2]  
  weight.in.kg <- weight / 2.2  
  height.in.m <- height * 0.0254  
  bmi <- weight.in.kg / (height.in.m ^ 2)  
  list(first.name=first.name, last.name=last.name, weight=weight.in.kg,  
    height=height.in.m,  
    bmi=bmi)  
}
```

91

## Try out function

```
createPatientRecord("Michael Smith", 185, 12 * 6 + 1)
```

```
$first.name  
[1] "Michael"  
  
$last.name  
[1] "Smith"  
  
$weight  
[1] 84.09091  
  
$height  
[1] 1.8542  
  
$bmi  
[1] 24.45884
```

92

# Another example

- Calculate mean, median and standard deviation

```
threeNumberSummary <- function(x) {  
  c(mean=mean(x), median=median(x), sd=sd(x))  
}  
x <- rnorm(100, mean=5, sd=2) # Vector of 100 normals with mean 5 and sd 2  
threeNumberSummary(x)
```

```
mean      median      sd  
4.926875 5.050984 1.953703
```

93

# If-else statement

- Oftentimes we want our code to have different effects depending on the features of the input
- Example: Calculating a student's letter grade
  - If grade  $\geq 90$ , assign A
  - Otherwise, if grade  $\geq 80$ , assign B
  - Otherwise, if grade  $\geq 70$ , assign C
  - In all other cases, assign F
- To code this up, we use if-else statements

94

# If-else example

```
calculateLetterGrade <- function(x) {  
  if(x >= 90) {  
    grade <- "A"  
  } else if(x >= 80) {  
    grade <- "B"  
  } else if(x >= 70) {  
    grade <- "C"  
  } else {  
    grade <- "F"  
  }  
  grade  
}  
  
course.grades <- c(92, 78, 87, 91, 62)  
sapply(course.grades, FUN=calculateLetterGrade)
```

```
[1] "A" "C" "B" "A" "F"
```

95

# return()

- In the previous examples we specified the output simply by writing the output variable as the last line of the function
- More explicitly, we can use the `return( )` function

```
addOne <- function(x) {  
  return(x + 1)  
}  
  
addOne(12)
```

```
[1] 13
```

- We will generally avoid the `return( )` function, but you can use it if necessary or if it makes writing a particular function easier.

96

# Workshop 4 : Transform

For the first two problems we'll use the Cars93 data set from the MASS library.

```
library(MASS)
```

## 1. Manipulating data frames

Use the **transform()** and **log()** functions to create a new data frame called Cars93.log that has MPG.highway and MPG.city replaced with log(MPG.highway) and log(MPG.city).

## 2. Functions, lists, and if-else

- Write a function called **isPassingGrade** whose input x is a number, and which returns FALSE if x is lower than 50 and TRUE otherwise.
- Write a function called **sendMessage** whose input x is a number, and which prints **Congratulations** if **isPassingGrade(x)** is TRUE and prints **Oh no!** if **isPassingGrade(x)** is FALSE.

97

# Data Wrangling and Visualization

# Common Problem

- One of the most common problems you'll encounter when importing manually-entered data is inconsistent data types within columns
- For a simple example, let's look at TVhours column in a messy version of the survey data

```
survey.messy <- read.csv("survey_messy.csv", header=TRUE)  
survey.messy$TVhours
```

```
[1] 1           8h          4           5  
[5] 3           ~10         0           10  
[9] 15 (incl movies) 4           0           0  
[13] 10          4           10          2hours  
[17] 5           3           0           1  
[21] 2           7           3           10  
[25] 6.5          40          0           12  
[29] adfjalkj     3           0           ...  
16 Levels: ~10 0 1 10 12 15 (incl movies) 2 2hours 3 4 40 5 6.5 7 ... adfjalkj
```

99

## What are the problems?

```
str(survey.messy)
```

```
'data.frame': 31 obs. of 6 variables:  
 $ Program      : Factor w/ 3 levels "MISM","Other",...: 1 2 1 3 2 3 2 1 3 3 ...  
 $ PriorExp     : Factor w/ 3 levels "Extensive experience",...: 3 3 1 2 2 3 3 3  
 2 1 ...  
 $ Rexperience   : Factor w/ 3 levels "Basic competence",...: 3 1 1 2 3 3 3 3 3 3 ...  
 ...  
 $ OperatingSystem: Factor w/ 2 levels "Mac OS X","Windows": 2 2 2 2 2 1 1 2 1 2  
 ...  
 $ TVhours       : Factor w/ 16 levels "~10","0","1",...: 3 15 10 12 9 1 2 4 6 10  
 ...  
 $ Editor        : Factor w/ 1 level "Microsoft Word": 1 1 1 1 1 1 1 1 1 1 ...
```

- Several of the entries have non-numeric values in them (they contain strings)
- As a result, TVhours is being imported as factor

100

# Fix the type

```
as.character(tv.hours.messy)[1:30]
```

```
[1] "1"          "8h"         "4"  
[4] "5"          "3"           "~10"  
[7] "0"          "10"          "15 (incl movies)"  
[10] "4"          "0"           "0"  
[13] "10"         "4"           "10"  
[16] "2hours"    "5"           "3"  
[19] "0"          "1"           "2"  
[22] "7"          "3"           "10"  
[25] "6.5"        "40"          "0"  
[28] "12"         "adfjalkj"   "3"
```

```
as.numeric(as.character(tv.hours.messy))[1:30]
```

```
[1] 1.0      NA     4.0    5.0    3.0    NA     0.0   10.0   NA    4.0    0.0    0.0   10.0   4.0  
[15] 10.0     NA    5.0    3.0    0.0    1.0    2.0    7.0    3.0   10.0   6.5    40.0   0.0   12.0  
[29] NA       3.0
```

```
typeof(as.numeric(as.character(tv.hours.messy))) # Success!! (Almost...)
```

```
[1] "double"
```

101

# A small improvement

- All the corrupted cells now appear as NA, which is R's missing indicator
- We can do a little better by cleaning up the vector once we get it to character form

```
tv.hours.strings <- as.character(tv.hours.messy)  
tv.hours.strings
```

```
[1] "1"          "8h"         "4"  
[4] "5"          "3"           "~10"  
[7] "0"          "10"          "15 (incl movies)"  
[10] "4"          "0"           "0"  
[13] "10"         "4"           "10"  
[16] "2hours"    "5"           "3"  
[19] "0"          "1"           "2"  
[22] "7"          "3"           "10"  
[25] "6.5"        "40"          "0"  
[28] "12"         "adfjalkj"   "3"  
[31] "0"
```

102

# Deleting non-numeric (or .) characters

```
tv.hours.strings
```

```
[1] "1"                 "8h"                "4"  
[4] "5"                 "3"                  "~10"  
[7] "0"                 "10"               "15 (incl movies)"  
[10] "4"                "0"                  "0"  
[13] "10"               "4"                  "10"  
[16] "2hours"            "5"                  "3"  
[19] "0"                 "1"                  "2"  
[22] "7"                 "3"                  "10"  
[25] "6.5"               "40"                "0"  
[28] "12"                "adfjalkj"           "3"  
[31] "0"
```

```
# Use gsub() to replace everything except digits and '.' with a blank ""  
gsub("[^0-9.]", "", tv.hours.strings)
```

```
[1] "1"    "8"    "4"    "5"    "3"    "10"   "0"    "10"   "15"   "4"    "0"  
[12] "0"   "10"   "4"    "10"   "2"    "5"    "3"    "0"    "1"    "2"    "7"  
[23] "3"   "10"   "6.5"  "40"   "0"    "12"   ""     "3"    "0"    "
```

103

## Redo

```
tv.hours.messy
```

```
[1] 1                 8h                4                 5  
[5] 3                 ~10               0                 10  
[9] 15 (incl movies) 4                 0                 0  
[13] 10                4                 10                2hours  
[17] 5                 3                 0                 1  
[21] 2                 7                 3                 10  
[25] 6.5               40                0                 12  
[29] adfjalkj           3                 0  
16 Levels: -10 0 1 10 12 15 (incl movies) 2 2hours 3 4 40 5 6.5 7 ... adfjalkj
```

```
tv.hours.clean <- as.numeric(gsub("[^0-9.]", "", tv.hours.strings))  
tv.hours.clean
```

```
[1] 1.0  8.0  4.0  5.0  3.0 10.0  0.0 10.0 15.0  4.0  0.0  0.0 10.0  4.0  
[15] 10.0 2.0  5.0  3.0  0.0  1.0  2.0  7.0  3.0 10.0  6.5 40.0  0.0 12.0  
[29] NA   3.0  0.0
```

104

# Another approach

- We can also handle this problem by setting `stringsAsFactors = FALSE` when importing our data.

```
survey.messy <- read.csv("survey_messy.csv", header=TRUE, stringsAsFactors=FALSE)
str(survey.messy)
```

```
'data.frame': 31 obs. of 6 variables:
 $ Program      : chr "MISM" "Other" "MISM" "PPM" ...
 $ PriorExp     : chr "Some experience" "Some experience" "Extensive
experience" "Never programmed before" ...
 $ Rexperience   : chr "Never used" "Basic competence" "Basic competence"
"Installed on machine" ...
 $ OperatingSystem: chr "Windows" "Windows" "Windows" "Windows" ...
 $ TVhours       : chr "1" "8h" "4" "5" ...
 $ Editor        : chr "Microsoft Word" "Microsoft Word" "Microsoft Word"
"Microsoft Word" ...
```

- Now everything is a character instead of a factor

105

# One-line Cleanup

- Let's clean up the `TVhours` column and cast it to numeric all in one command

```
survey <- transform(survey.messy, TVhours = as.numeric(gsub("[^0-9.]", "", 
TVhours)))
str(survey)
```

```
'data.frame': 31 obs. of 6 variables:
 $ Program      : chr "MISM" "Other" "MISM" "PPM" ...
 $ PriorExp     : chr "Some experience" "Some experience" "Extensive
experience" "Never programmed before" ...
 $ Rexperience   : chr "Never used" "Basic competence" "Basic competence"
"Installed on machine" ...
 $ OperatingSystem: chr "Windows" "Windows" "Windows" "Windows" ...
 $ TVhours       : num 1 8 4 5 3 10 0 10 15 4 ...
 $ Editor        : chr "Microsoft Word" "Microsoft Word" "Microsoft Word"
"Microsoft Word" ...
```

106

# What about other character column?

```
table(survey[["Program"]])
```

MISM	Other	PPM
10	6	15

```
table(as.factor(survey[["Program"]]))
```

MISM	Other	PPM
10	6	15

- Having factors coded as characters may be OK for many parts of our analysis

107

## Let's fix it

```
# Figure out which columns are coded as characters  
chr.indexes <- sapply(survey, FUN = is.character)  
chr.indexes
```

Program	PriorExp	Rexperience	OperatingSystem
TRUE	TRUE	TRUE	TRUE
TVhours	Editor		
FALSE	TRUE		

```
# Re-code all of the character columns to factors  
survey[chr.indexes] <- lapply(survey[chr.indexes], FUN = as.factor)
```

108

# Here is the outcome

```
str(survey)
```

```
'data.frame': 31 obs. of 6 variables:  
 $ Program      : Factor w/ 3 levels "MISM","Other",...: 1 2 1 3 2 3 2 1 3 3 ...  
 $ PriorExp     : Factor w/ 3 levels "Extensive experience",...: 3 3 1 2 2 3 3 3  
 2 1 ...  
 $ Rexperience   : Factor w/ 3 levels "Basic competence",...: 3 1 1 2 3 3 3 3 3 3  
 ...  
 $ OperatingSystem: Factor w/ 2 levels "Mac OS X","Windows": 2 2 2 2 2 1 1 2 1 2  
 ...  
 $ TVhours       : num  1 8 4 5 3 10 0 10 15 4 ...  
 $ Editor        : Factor w/ 1 level "Microsoft Word": 1 1 1 1 1 1 1 1 1 1 ...
```

- Success!

109

# Another example

- When data is entered manually, misspellings and case changes are very common
- E.g., a column showing life support mechanism may look like,

```
life.support <- as.factor(c("dialysis", "Ventilation", "Dialysis", "dialysis",  
 "none", "None", "nnone", "dyalysis", "dialysis", "ventilation", "none"))  
summary(life.support)
```

dialysis	Dialysis	dyalysis	nnone	none	None
3	1	1	1	2	1
ventilation	Ventilation				
1	1				

110



```
summary(life.support)
```

dialysis	Dialysis	dyalysis	nnone	none	None
3	1	1	1	2	1
ventilation	Ventilation				
1	1				

- This factor has 8 levels even though it should have 3 (dialysis, ventilation, none)
- We can fix many of the typos by running spellcheck in Excel before importing data, or by changing the values on a case-by-case basis later

111

## [l/s/t]apply() functions?

- These are all efficient ways of applying a function to margins of an array or elements of a list
- Before we talk about the details of `apply()` and its relatives, we should first understand loops
- **loops** are ways of iterating over data
- The `apply()` functions can be thought of as good *alternatives* to loops

112

# Loop

```
for(i in 1:4) {  
  print(i)  
}
```

```
[1] 1  
[1] 2  
[1] 3  
[1] 4
```

```
phrase <- "Good Night, "  
for(word in c("and", "Good", "Luck")) {  
  phrase <- paste(phrase, word)  
  print(phrase)  
}
```

```
[1] "Good Night, and"  
[1] "Good Night, and Good"  
[1] "Good Night, and Good Luck"
```

113

# Loop Syntax

A **for loop** executes a chunk of code for every value of an **index variable** in an **index set**

- The basic syntax takes the form

```
for(index.variable in index.set) {  
  code to be repeated at every value of index.variable  
}
```

- The index set is often a vector of integers, but can be more general

114

# Example

```
index.set <- rnorm(4) # 4 random standard normal variables  
index.set
```

```
[1] -1.9912818  0.7119512  0.3999575  0.5872188
```

```
for(i in index.set) {  
  print(abs(i))  
}
```

```
[1] 1.991282  
[1] 0.7119512  
[1] 0.3999575  
[1] 0.5872188
```

115

# Example

```
index.set <- list(name="Michael", weight=185, is.male=TRUE) # a list  
for(i in index.set) {  
  print(c(i, typeof(i)))  
}
```

```
[1] "Michael"    "character"  
[1] "185"        "double"  
[1] "TRUE"        "logical"
```

116

# Example

```
fake.data <- matrix(rnorm(500), ncol=5) # create fake 100 x 5 data set
head(fake.data,2) # print first two rows
```

```
[,1]      [,2]      [,3]      [,4]      [,5]
[1,] 0.9551196 -0.3985364 -0.2338078 -1.018674 -0.684313
[2,] 0.6483192 -0.5487889  0.2091660  2.016935 -1.030668
```

```
col.sums <- numeric(ncol(fake.data)) # variable to store running column sums
for(i in 1:nrow(fake.data)) {
  col.sums <- col.sums + fake.data[i,] # add ith observation to the sum
}
print(col.sums)
```

```
[1] -12.808804 -11.552563  3.211727  4.049877 -9.901475
```

```
colSums(fake.data) # A better approach (see also colMeans())
```

```
[1] -12.808804 -11.552563  3.211727  4.049877 -9.901475
```

117

# While loop

- **while loops** repeat a chunk of code while the specified condition remains true

```
day <- 1
num.days <- 365
while(day <= num.days) {
  day <- day + 1
}
```

- We won't really be using while loops in this class

118

# Various app() function

Command	Description
apply(X, MARGIN, FUN)	Obtain a vector/array/list by applying FUN along the specified MARGIN of an array or matrix X
lapply(X, FUN)	Obtain a list by applying FUN to the elements of a list x
sapply(X, FUN)	Simplified version of lapply. Returns a vector/array instead of list.
tapply(X, INDEX, FUN)	Obtain a table by applying FUN to each combination of the factors given in INDEX

- These functions are (good!) alternatives to loops
- They are typically *more efficient* than loops (often run considerably faster on large data sets)
- Take practice to get used to, but make analysis easier to debug and less prone to error when used effectively

119

## Example: apply()

```
colMeans(fake.data)  
[1] -0.12808804 -0.11552563  0.03211727  0.04049877 -0.09901475  
  
apply(fake.data, MARGIN=2, FUN=mean) # MARGIN = 1 for rows, 2 for columns  
[1] -0.12808804 -0.11552563  0.03211727  0.04049877 -0.09901475  
  
# Function that calculates proportion of vector indexes that are > 0  
propPositive <- function(x) mean(x > 0)  
apply(fake.data, MARGIN=2, FUN=propPositive)  
[1] 0.51 0.47 0.51 0.51 0.46
```

120

# Example: lapply(), sapply()

```
lapply(survey, is.factor) # Returns a list
```

```
$Program  
[1] TRUE  
  
$PriorExp  
[1] TRUE  
  
$Rexperience  
[1] TRUE  
  
$OperatingSystem  
[1] TRUE  
  
$TVhours  
[1] FALSE  
  
$Editor  
[1] TRUE
```

```
sapply(survey, FUN = is.factor) # Returns a vector with named elements
```

Program	PriorExp	Rexperience	OperatingSystem
TRUE	TRUE	TRUE	TRUE
TVhours	Editor		
FALSE	TRUE		

121

# Example: apply, lapply, sapply

```
apply(cars, 2, FUN=mean) # Data frames are arrays
```

```
speed dist  
15.40 42.98
```

```
lapply(cars, FUN=mean) # Data frames are also lists
```

```
$speed  
[1] 15.4  
  
$dist  
[1] 42.98
```

```
sapply(cars, FUN=mean) # sapply() is just simplified lapply()
```

```
speed dist  
15.40 42.98
```

122

# tapply()

- Think of tapply() as a generalized form of the table() function

```
library(MASS)
# Get a count table, data broken down by Origin and DriveTrain
table(Cars93$Origin, Cars93$DriveTrain)
```

	4WD	Front	Rear
USA	5	34	9
non-USA	5	33	7

```
# Calculate average MPG.City, broken down by Origin and Drivetrain
tapply(Cars93$MPG.city, INDEX = Cars93[c("Origin", "DriveTrain")], FUN=mean)
```

	DriveTrain		
Origin	4WD	Front	Rear
USA	17.6	22.14706	18.33333
non-USA	23.4	24.93939	19.14286

123

## Example: tapply()

- Let's get the average horsepower by car Origin and Type

```
tapply(Cars93[["Horsepower"]], INDEX = Cars93[c("Origin", "Type")], FUN=mean)
```

	Type						
Origin	Compact	Large	Midsize	Small	Sporty	Van	
USA	117.4286	179.4545	153.5000	89.42857	166.5000	158.40	
non-USA	141.5556	NA	189.4167	91.78571	151.6667	138.25	

- What's that NA doing there?

```
any(Cars93$Origin == "non-USA" & Cars93$type == "Large")
```

```
[1] FALSE
```

- None of the non-USA manufacturers produced Large cars!

124

# with()

- Thus far we've repeatedly typed out the data frame name when referencing its columns
- This is because the data variables don't exist in our working environment
- Using **with**(*data*, *expr*) lets us specify that the code in *expr* should be evaluated in an environment that contains the elements of *data* as variables

```
with(Cars93, table(Origin, Type))
```

Origin	Type					
	Compact	Large	Midsize	Small	Sporty	Van
USA	7	11	10	7	8	5
non-USA	9	0	12	14	6	4

125

## Example: with()

```
any(Cars93$Origin == "non-USA" & Cars93>Type == "Large")
```

```
[1] FALSE
```

```
with(Cars93, any(Origin == "non-USA" & Type == "Large")) # Same effect!
```

```
[1] FALSE
```

```
with(Cars93, tapply(Horsepower, INDEX = list(Origin, Type), FUN=mean))
```

	Compact	Large	Midsize	Small	Sporty	Van
USA	117.4286	179.4545	153.5000	89.42857	166.5000	158.40
non-USA	141.5556	NA	189.4167	91.78571	151.6667	138.25

- Using **with()** makes code simpler, easier to read, and easier to debug

126

# Workshop 6: Loop

## Loop practice

- (a) Write a function called calculateRowMeans that uses a for loop to calculate the row means of a matrix x.
- (b) Try out your function on the random matrix fake.data defined below.
- (b) Use the apply() function to calculate the row means of the matrix fake.data
- (c) Compare this to the output of the rowMeans() function to check that your calculation is correct.

127

# Data Summary and Standard Graphic

# Let's get started

- We're going to start by operating on the `birthwt` dataset from the MASS library
- Let's get it loaded and see what we're working with

```
library(MASS)
str(birthwt)
```

```
'data.frame': 189 obs. of 10 variables:
 $ low   : int  0 0 0 0 0 0 0 0 0 ...
 $ age   : int  19 33 20 21 18 21 22 17 29 26 ...
 $ lwt   : int  182 155 105 108 107 124 118 103 123 113 ...
 $ race  : int  2 3 1 1 1 3 1 3 1 1 ...
 $ smoke : int  0 0 1 1 1 0 0 0 1 1 ...
 $ ptl   : int  0 0 0 0 0 0 0 0 0 0 ...
 $ ht    : int  0 0 0 0 0 0 0 0 0 0 ...
 $ ui   : int  1 0 0 1 1 0 0 0 0 0 ...
 $ ftv   : int  0 3 1 2 0 0 1 1 1 0 ...
 $ bwt   : int  2523 2551 2557 2594 2600 2622 2637 2637 2663 2665 ...
```

129

## Renaming columns

- The dataset doesn't come with very descriptive variable names
- Let's get better column names (use `help(birthwt)` to understand the variables and come up with better names)

```
colnames(birthwt)
```

```
[1] "low"     "age"      "lwt"      "race"     "smoke"    "ptl"      "ht"       "ui"
[9] "ftv"     "bwt"
```

```
# The default names are not very descriptive
```

```
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",
                      "race", "mother.smokes", "previous.prem.labor", "hypertension",
                      "uterine.irr",
                      "physician.visits", "birthwt.grams")
```

```
# Better names!
```

130

# Renaming the variables

- The dataset doesn't come with very descriptive variable names
- Let's get better column names (use `help(birthwt)` to understand the variables and come up with better names)

```
colnames(birthwt)  
[1] "low"     "age"      "lwt"      "race"     "smoke"    "ptl"      "ht"       "ui"  
[9] "ftv"     "bwt"  
  
# The default names are not very descriptive  
  
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",  
  "race", "mother.smokes", "previous.prem.labor", "hypertension",  
  "uterine.irr",  
  "physician.visits", "birthwt.grams")  
  
# Better names!
```

131

# Renaming the factors

- All the factors are currently represented as integers
- Let's use the `transform()` and `mapvalues()` functions to convert variables to factors and give the factors more meaningful levels

```
library(plyr)  
birthwt <- transform(birthwt,  
  race = as.factor(mapvalues(race, c(1, 2, 3),  
    c("white", "black", "other"))),  
  mother.smokes = as.factor(mapvalues(mother.smokes,  
    c(0,1), c("no", "yes"))),  
  hypertension = as.factor(mapvalues(hypertension,  
    c(0,1), c("no", "yes"))),  
  uterine.irr = as.factor(mapvalues(uterine.irr,  
    c(0,1), c("no", "yes"))),  
  birthwt.below.2500 = as.factor(mapvalues(birthwt.below.2500,  
    c(0,1), c("no", "yes"))))
```

132

# Summary of the data

- Now that things are coded correctly, we can look at an overall summary

```
summary(birthwt)
```

```
birthwt.below.2500    mother.age      mother.weight      race
no :130                Min. :14.00     Min.   : 80.0       black:26
yes: 59               1st Qu.:19.00    1st Qu.:110.0     other:67
                           Median :23.00     Median :121.0     white:96
                           Mean   :23.24     Mean   :129.8
                           3rd Qu.:26.00    3rd Qu.:140.0
                           Max.  :45.00     Max.  :250.0
mother.smokes previous.prem.labor hypertension uterine.irr
no :115                Min. :0.0000    no :177        no :161
yes: 74                 1st Qu.:0.0000  yes: 12       yes: 28
                           Median :0.0000
                           Mean  :0.1958
                           3rd Qu.:0.0000
                           Max.  :3.0000
physician.visits birthwt.grams
Min.  :0.0000  Min.   : 709
1st Qu.:0.0000 1st Qu.:2414
Median :0.0000  Median :2977
Mean   :0.7937  Mean   :2945
3rd Qu.:1.0000 3rd Qu.:3487
Max.   :6.0000  Max.   :4990
```

133

# A simple table

- Let's use the `tapply()` function to see what the average birthweight looks like when broken down by race and smoking status

```
with(birthwt, tapply(birthwt.grams, INDEX = list(race, mother.smokes), FUN =
mean))
```

	no	yes
black	2854.500	2504.000
other	2815.782	2757.167
white	3428.750	2826.846

134

# Output table

- Let's use the header `{r, results='asis'}`, along with the `kable()` function from the `knitr` library

```
library(knitr)
kable(with(birthwt, tapply(birthwt.grams, INDEX = list(race, mother.smokes), FUN = mean)), format = "markdown")
```

	no	yes
black	2854.500	2504.000
other	2815.782	2757.167
white	3428.750	2826.846

- `kable()` outputs the table in a way that Markdown can read and nicely display
- Note: changing the CSS changes the table appearance

135

# aggregate() function

- Let's first recall what `tapply()` does
- Command: `tapply(X, INDEX, FUN)`
  - Applies `FUN` to `X` grouped by factors in `INDEX`
- `aggregate()` performs a similar operation, but presents the results in a form that is at times more convenient
- There are many ways to call the `aggregate()` function
- Analog of `tapply` call: `aggregate(X, by, FUN)`
  - Here, `by` is exactly like `INDEX`

136

# Example: tapply vs aggregate

```
library(MASS)
with(birthwt, tapply(birthwt.grams, INDEX = list(race, mother.smokes), FUN =
mean)) # tapply
```

```
      no      yes
black 2854.500 2504.000
other 2815.782 2757.167
white 3428.750 2826.846
```

```
with(birthwt, aggregate(birthwt.grams, by = list(race, mother.smokes), FUN =
mean)) # aggregate
```

```
  Group.1 Group.2      x
1   black      no 2854.500
2   other      no 2815.782
3   white      no 3428.750
4   black     yes 2504.000
5   other     yes 2757.167
6   white     yes 2826.846
```

137

## Different syntax

- Here's a convenient alternative way to call aggregate
- It uses the R formula syntax, which we'll learn more about when we discuss regression

```
aggregate(birthwt.grams ~ race + mother.smokes, FUN=mean, data=birthwt)
```

```
  race mother.smokes birthwt.grams
1 black      no 2854.500
2 other      no 2815.782
3 white      no 3428.750
4 black     yes 2504.000
5 other     yes 2757.167
6 white     yes 2826.846
```

- We'll see later that aggregate output can be more convenient for plotting

138



```
weight.smoke.tbl <- with(birthwt, table(birthwt.below.2500, mother.smokes))  
weight.smoke.tbl
```

		mother.smokes	
		no	yes
birthwt.below.2500	no	86	44
	yes	29	30

- The odds of low bwt among non-smoking mothers is

```
or.smoke.bwt <- (weight.smoke.tbl[2,2] / weight.smoke.tbl[1,2]) /  
(weight.smoke.tbl[2,1] / weight.smoke.tbl[1,1])  
or.smoke.bwt
```

```
[1] 2.021944
```

- So the odds of low birth weight are 2 times higher when the mother smokes

139



- Is the mother's age correlated with birth weight?

```
with(birthwt, cor(birthwt.grams, mother.age)) # Calculate correlation  
  
[1] 0.09031781
```

- Does this change when we account for smoking status?

```
with(birthwt, cor(birthwt.grams[mother.smokes == "yes"],  
mother.age[mother.smokes == "yes"]))  
  
[1] -0.1441649  
  
with(birthwt, cor(birthwt.grams[mother.smokes == "no"], mother.age[mother.smokes  
== "no"]))  
  
[1] 0.2014558
```

140

# by() function

- Think of the `by(data, INDICES, FUN)` function as a `tapply()` function that operates on data frames instead of just vectors
- When using `tapply(X, INDEX, FUN)`, `X` is generally a numeric vector
- To calculate correlations, we need to allow `x` to be a data frame or matrix

```
by(data = birthwt[c("birthwt.grams", "mother.age")],  
   INDICES = birthwt["mother.smokes"],  
   FUN = function(x) {cor(x[,1], x[,2])})
```

```
mother.smokes: no  
[1] 0.2014558
```

```
-----  
mother.smokes: yes  
[1] -0.1441649
```

141

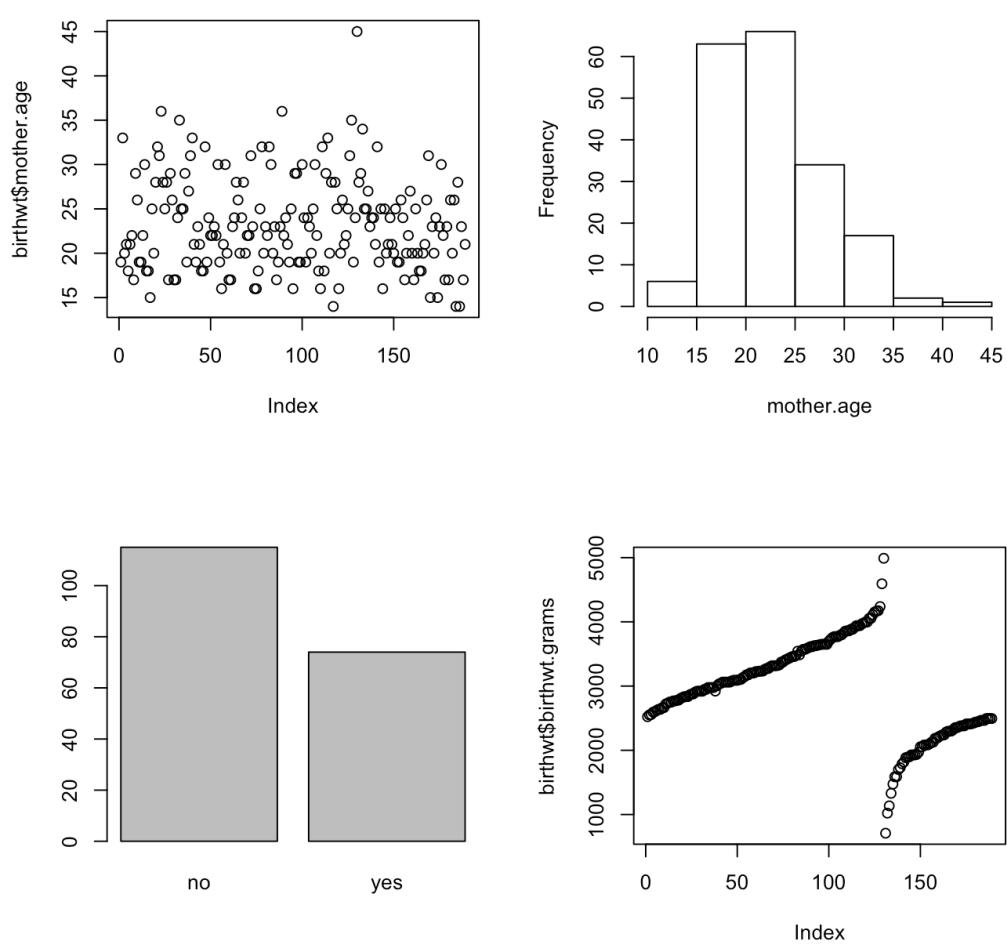
# Basic single plot

Let's continue with the `birthwt` data from the `MASS` library.

Here are some basic single-variable plots.

```
par(mfrow = c(2,2)) # Display plots in a single 2 x 2 figure  
plot(birthwt$mother.age)  
with(birthwt, hist(mother.age))  
plot(birthwt$mother.smokes)  
plot(birthwt$birthwt.grams)
```

142

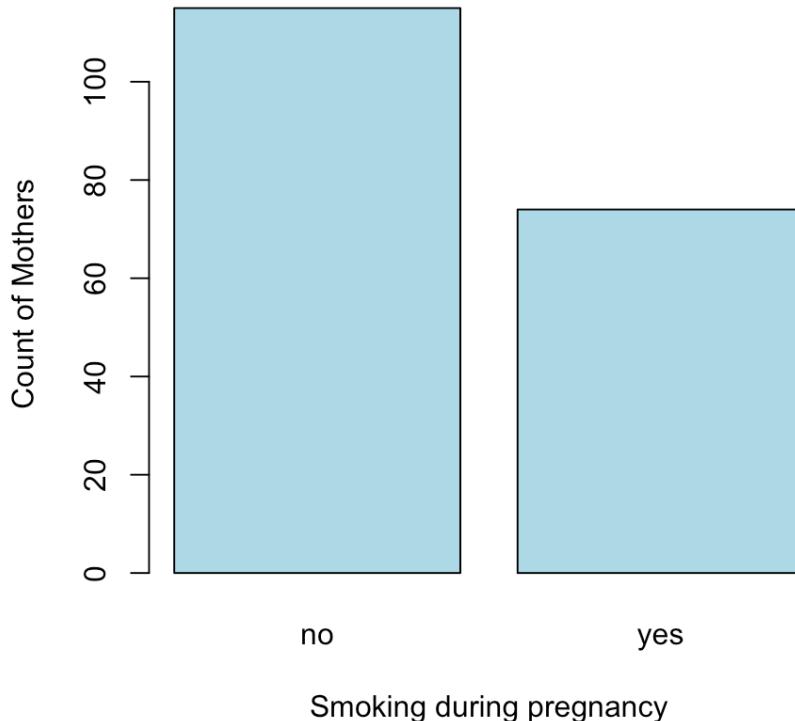


## Another example

Let's add more information to the smoking bar plot, and also change the color by setting the `col` option.

```
plot(birthwt$mother.smokes,
      main = "Mothers Who Smoked In Pregnancy",
      xlab = "Smoking during pregnancy",
      ylab = "Count of Mothers",
      col = 'lightblue')
```

## Mothers Who Smoked In Pregnancy



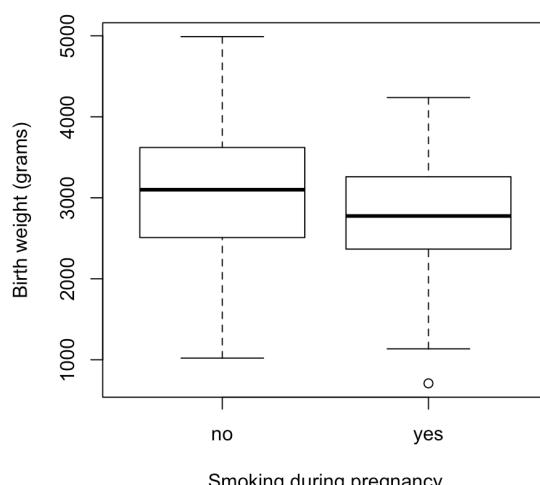
145

## Plots with several variables

If we call `plot(x, y, ...)` with `x` a factor and `y` numeric, R will produce boxplots of `y` at every level of `x`.

```
with(birthwt, plot(mother.smokes, birthwt.grams,
                    main = "Birth Weight by Smoking Status",
                    xlab = "Smoking during pregnancy",
                    ylab = "Birth weight (grams)"))
```

Birth Weight by Smoking Status



146

# More on ggplot2 demo

147

## Workshop 7: Plot data

Plotting the Cars93 data

- (a) Use qplot to create a scatterplot with Price on the y-axis and EngineSize on the x-axis.
- (b) Repeat part (a) using the ggplot function and geom\_point() layer.
- (c) Repeat part (c), but this time specifying that the color mapping should depend on Type and the shape mapping should depend on DriveTrain.

148



# Other Important Packages



## To load data

**RODBC, RMySQL, RPostgreSQL, RSQLite** - If you'd like to read in data from a database, these packages are a good place to start. Choose the package that fits your type of database.

**XLConnect, xlsx** - These packages help you read and write Microsoft Excel files from R. You can also just export your spreadsheets from Excel as .csv's.

**foreign** - Want to read a SAS data set into R? Or an SPSS data set? Foreign provides functions that help you load data files from other programs into R.

R can handle plain text files – no package required. Just use the functions `read.csv`, `read.table`, and `read.fwf..`

# xlsx

```
df <- read.xlsx("<name and extension of your file>",
                 sheetIndex = 1)
```

**Note** that it is necessary to add a sheet name or a sheet index to this function. In the example above, the first sheet of the Excel file was assigned. If you have a bigger data set, you might get better performance when using the `read.xlsx2()` function:

```
df <- read.xlsx2("<name and extension of your file>",
                  sheetIndex = 1,
                  startRow=2,
                  colIndex = 2)
```

151

```
write.xlsx(df,
           "df.xlsx",
           sheetName="Data Frame")
```

The function requires you first to specify what data frame you want to export. In the second argument, you specify the name of the file that you are outputting. If, however, you want to write the data frame to a file that already exists, you can execute the following command:

```
write.xlsx(df,
           "<name and extension of your existing file>",
           sheetName="Data Frame"
           append=TRUE)
```

152

# To manipulate data

**dplyr** - Essential shortcuts for subsetting, summarizing, rearranging, and joining together data sets. dplyr is our go to package for fast data manipulation.

**tidyverse** - Tools for changing the layout of your data sets. Use the gather and spread functions to convert your data into the tidy format, the layout R likes best.

**stringr** - Easy to learn tools for regular expressions and character strings.

**lubridate** - Tools that make working with dates and times easier.

153

## dplyr

```
install.packages("dplyr")
```

```
library(dplyr)
```

dplyr verbs	Description
<code>select()</code>	select columns
<code>filter()</code>	filter rows
<code>arrange()</code>	re-order or arrange rows
<code>mutate()</code>	create new columns
<code>summarise()</code>	summarise values
<code>group_by()</code>	allows for group operations in the “split-apply-combine” concept

154

# select

Select a set of columns: the name and the sleep\_total columns.

```
sleepData <- select(msleep, name, sleep_total)  
head(sleepData)
```

```
##           name sleep_total  
## 1      Cheetah     12.1  
## 2    Owl monkey    17.0  
## 3 Mountain beaver 14.4  
## 4 Greater short-tailed shrew 14.9  
## 5          Cow     4.0  
## 6 Three-toed sloth 14.4
```

155

# Subtraction (-)

To select all the columns *except* a specific column, use the “-“ (subtraction) operator (also known as negative indexing)

```
head(select(msleep, -name))
```

```
##           genus vore      order conservation sleep_total sleep_rem  
## 1  Acinonyx carni Carnivora         lc     12.1      NA  
## 2      Aotus omni Primates        <NA>    17.0      1.8  
## 3   Aplodontia herbi Rodentia        nt     14.4      2.4  
## 4      Blarina omni Soricomorpha      lc     14.9      2.3  
## 5      Bos herbi Artiodactyla domesticated 4.0      0.7  
## 6   Bradypus herbi Pilosa        <NA>    14.4      2.2  
##           sleep_cycle awake brainwt bodywt  
## 1             NA  11.9      NA 50.000  
## 2             NA   7.0 0.01550  0.480  
## 3             NA   9.6      NA  1.350  
## 4  0.1333333  9.1 0.00029  0.019  
## 5  0.6666667 20.0 0.42300 600.000  
## 6  0.7666667   9.6      NA  3.850
```

156

# colon (:)

To select a range of columns by name, use the “:” (colon) operator

```
head(select(msleep, name:order))
```

```
##          name      genus  vore      order
## 1       Cheetah  Acinonyx carni  Carnivora
## 2     Owl monkey     Aotus  omni   Primates
## 3 Mountain beaver  Aplodontia herbi Rodentia
## 4 Greater short-tailed shrew    Blarina  omni Soricomorpha
## 5            Cow        Bos herbi Artiodactyla
## 6 Three-toed sloth   Bradypus herbi   Pilosa
```

157



To select all columns that start with the character string “sl”, use the function `starts_with()`

```
head(select(msleep, starts_with("sl")))
```

```
##   sleep_total sleep_rem sleep_cycle
## 1      12.1        NA        NA
## 2      17.0        1.8        NA
## 3      14.4        2.4        NA
## 4      14.9        2.3  0.1333333
## 5       4.0        0.7  0.6666667
## 6      14.4        2.2  0.7666667
```

Some additional options to select columns based on a specific criteria include

1. `ends_with()` = Select columns that end with a character string
2. `contains()` = Select columns that contain a character string
3. `matches()` = Select columns that match a regular expression
4. `one_of()` = Select columns names that are from a group of names

158

# filter

Filter the rows for mammals that sleep a total of more than 16 hours.

```
filter(msleep, sleep_total >= 16)
```

```
##                 name      genus   vore      order conservation
## 1     Owl monkey     Aotus    omni Primates      <NA>
## 2 Long-nosed armadillo Dasypus   carni Cingulata       lc
## 3 North American Opossum Didelphis  omni Didelphimorphia   lc
## 4     Big brown bat Eptesicus insecti Chiroptera       lc
## 5 Thick-tailed opossum Lutreolina  carni Didelphimorphia   lc
## 6     Little brown bat Myotis insecti Chiroptera      <NA>
## 7     Giant armadillo Priodontes insecti Cingulata       en
## 8 Arctic ground squirrel Spermophilus  herbi Rodentia       lc
##   sleep_total sleep_rem sleep_cycle awake brainwt bodywt
## 1     17.0        1.8          NA    7.0  0.01550  0.480
## 2     17.4        3.1  0.3833333  6.6  0.01080  3.500
## 3     18.0        4.9  0.3333333  6.0  0.00630  1.700
## 4     19.7        3.9  0.1166667  4.3  0.00030  0.023
## 5     19.4        6.6          NA    4.6      NA  0.370
## 6     19.9        2.0  0.2000000  4.1  0.00025  0.010
## 7     18.1        6.1          NA    5.9  0.08100 60.000
## 8     16.6        NA          NA    7.4  0.00570  0.920
```

# filter with and

Filter the rows for mammals that sleep a total of more than 16 hours *and* have a body weight of greater than 1 kilogram.

```
filter(msleep, sleep_total >= 16, bodywt >= 1)
```

```
##                 name      genus   vore      order conservation
## 1 Long-nosed armadillo Dasypus   carni Cingulata       lc
## 2 North American Opossum Didelphis  omni Didelphimorphia   lc
## 3     Giant armadillo Priodontes insecti Cingulata       en
##   sleep_total sleep_rem sleep_cycle awake brainwt bodywt
## 1     17.4        3.1  0.3833333  6.6  0.0108     3.5
## 2     18.0        4.9  0.3333333  6.0  0.0063     1.7
## 3     18.1        6.1          NA    5.9  0.0810    60.0
```

# Pipe (%>%)

```
head(select(msleep, name, sleep_total))
```

```
##           name sleep_total
## 1          Cheetah     12.1
## 2      Owl monkey    17.0
## 3 Mountain beaver   14.4
## 4 Greater short-tailed shrew 14.9
## 5          Cow       4.0
## 6 Three-toed sloth  14.4
```

161

```
msleep %>%
  select(name, sleep_total) %>%
  head
```

```
##           name sleep_total
## 1          Cheetah     12.1
## 2      Owl monkey    17.0
## 3 Mountain beaver   14.4
## 4 Greater short-tailed shrew 14.9
## 5          Cow       4.0
## 6 Three-toed sloth  14.4
```

162

# sqldf

```
library(sqldf)
```

```
setwd()
crashes <- read.csv("crashes.csv")
roads <- read.csv("roads.csv")
head(crashes)
```

```
##   Year      Road N_Crashes Volume
## 1 1991 Interstate 65      25 40000
## 2 1992 Interstate 65      37 41000
## 3 1993 Interstate 65      45 45000
## 4 1994 Interstate 65      46 45600
## 5 1995 Interstate 65      46 49000
## 6 1996 Interstate 65      59 51000
```

```
tail(crashes)
```

```
##   Year      Road N_Crashes Volume
## 105 2007 Interstate 275      32 21900
## 106 2008 Interstate 275      21 21850
## 107 2009 Interstate 275      25 22100
## 108 2010 Interstate 275      24 21500
## 109 2011 Interstate 275      23 20300
## 110 2012 Interstate 275      22 21200
```

163

```
print(roads)
```

```
##      Road      District Length
## 1 Interstate 65 Greenfield    262
## 2 Interstate 70 Vincennes   156
## 3 US-36 Crawfordsville 139
## 4 US-40 Greenfield    150
## 5 US-52 Crawfordsville 172
```

164

# Outer join

```
join_string <- "select
  crashes.*
, roads.District
, roads.Length
from crashes
  left join roads
    on crashes.Road = roads.Road"
```

```
crashes_join_roads <- sqldf(join_string,stringsAsFactors = FALSE)
```

```
## Loading required package: tcltk
```

```
head(crashes_join_roads)
```

```
##   Year      Road N_Crashes Volume District Length
## 1 1991 Interstate 65       25 40000 Greenfield   262
## 2 1992 Interstate 65       37 41000 Greenfield   262
## 3 1993 Interstate 65       45 45000 Greenfield   262
## 4 1994 Interstate 65       46 45600 Greenfield   262
## 5 1995 Interstate 65       46 49000 Greenfield   262
## 6 1996 Interstate 65       59 51000 Greenfield   262
```

165

# Inner join

By using an inner join, only matching rows will be kept.

```
join_string2 <- "select
  crashes.*
, roads.District
, roads.Length
from crashes
  inner join roads
    on crashes.Road = roads.Road"
```

```
crashes_join_roads2 <- sqldf(join_string2, stringsAsFactors = FALSE)
head(crashes_join_roads2)
```

```
##   Year      Road N_Crashes Volume District Length
## 1 1991 Interstate 65       25 40000 Greenfield   262
## 2 1992 Interstate 65       37 41000 Greenfield   262
## 3 1993 Interstate 65       45 45000 Greenfield   262
## 4 1994 Interstate 65       46 45600 Greenfield   262
## 5 1995 Interstate 65       46 49000 Greenfield   262
## 6 1996 Interstate 65       59 51000 Greenfield   262
```

166



```
tail(crashes_join_roads2)
```

```
##   Year Road N_Crashes Volume      District Length
## 83 2007 US-36        49 24000 Crawfordsville    139
## 84 2008 US-36        52 24500 Crawfordsville    139
## 85 2009 US-36        55 24700 Crawfordsville    139
## 86 2010 US-36        35 23000 Crawfordsville    139
## 87 2011 US-36        33 21000 Crawfordsville    139
## 88 2012 US-36        31 20500 Crawfordsville    139
```

167

## Where

```
join_string2 <- "select
  crashes.*
  , roads.District
  , roads.Length
  from crashes
  inner join roads
  on crashes.Road = roads.Road
  where crashes.Road = 'US-40'"
crashes_join_roads4 <- sqldf(join_string2,stringsAsFactors = FALSE)
head(crashes_join_roads4)
```

```
##   Year Road N_Crashes Volume      District Length
## 1 1991 US-40        46 21000 Greenfield    150
## 2 1992 US-40       101 21500 Greenfield    150
## 3 1993 US-40        76 23000 Greenfield    150
## 4 1994 US-40        72 21000 Greenfield    150
## 5 1995 US-40        75 24000 Greenfield    150
## 6 1996 US-40       136 23500 Greenfield    150
```

```
tail(crashes_join_roads4)
```

```
##   Year Road N_Crashes Volume      District Length
## 17 2007 US-40        45 59500 Greenfield    150
## 18 2008 US-40        23 61000 Greenfield    150
## 19 2009 US-40        67 65000 Greenfield    150
## 20 2010 US-40       102 67000 Greenfield    150
## 21 2011 US-40        87 67500 Greenfield    150
## 22 2012 US-40        32 67500 Greenfield    150
```

168

# To visualize data

**ggplot2** - R's famous package for making beautiful graphics. ggplot2 lets you use the grammar of graphics to build layered, customizable plots.

**ggvis** - Interactive, web based graphics built with the grammar of graphics.

**rgl** - Interactive 3D visualizations with R

169

**htmlwidgets** - A fast way to build interactive (javascript based) visualizations with R. Packages that implement htmlwidgets include:

**leaflet** (maps)

**dygraphs** (time series)

**DT** (tables)

**diagrammeR** (diagrams)

**network3D** (network graphs)

**threeJS** (3D scatterplots and globes).

170



**googleVis** - Let's you use Google Chart tools to visualize data in R.  
Google Chart tools used to be called Gapminder, the graphing software  
Hans Rosling made famous in his TED talk.

171

## To model data

**car** - car's Anova function is popular for making type II and type III Anova tables.

**mgcv** - Generalized Additive Models

**lme4/nlme** - Linear and Non-linear mixed effects models

**randomForest** - Random forest methods from machine learning

**multcomp** - Tools for multiple comparison testing

**vcd** - Visualization tools and tests for categorical data

**glmnet** - Lasso and elastic-net regression methods with cross validation

**survival** - Tools for survival analysis

172

# To report results

**shiny** - Easily make interactive, web apps with R. A perfect way to explore data and share findings with non-programmers.

**R Markdown** - The perfect workflow for reproducible reporting. Write R code in your markdown reports. When you run render, R Markdown will replace the code with its results and then export your report as an HTML, pdf, or MS Word document, or a HTML or pdf slideshow. The result?

Automated reporting. R Markdown is integrated straight into RStudio.

**xtable** - The xtable function takes an R object (like a data frame) and returns the latex or HTML code you need to paste a pretty version of the

173

## Introduction to Shiny

- Open Sourced by RStudio November 2012
- Default widgets and settings make it easy to generate apps
- Don't need to know HTML, CSS and javascript to get started
- Twitter Bootstrap for default UI - looks good
- Web sockets for communication between client and server
- Reactive Programming model
- Works on Windows, Mac, Linux

174

# Introduction to Shiny

Ready to use shiny

- Install R from CRAN
- Useful to have Chrome, Firefox, Safari...
- Install Shiny using R command: `install.packages("shiny")`

175

## Simple Example

```
• library(shiny)  
• runExample("01_hello")
```

Hello Shiny!

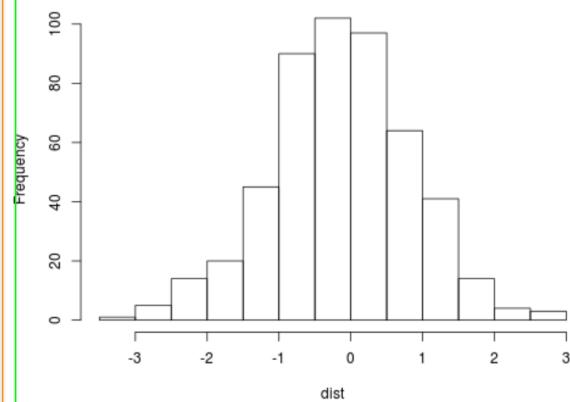
headerPanel() - Title

Number of observations:  
1            500            1,000

sidebarPanel() - Input

Histogram of dist

mainPanel() - Output



176

# ui.R :: Controls the look of the App

```
library(shiny)

# Define UI for application that draws a histogram
shinyUI(fluidPage(

  # Application title
  titlePanel("Hello Shiny!"),

  # Sidebar with a slider input for the number of bins
  sidebarLayout(
    sidebarPanel(
      sliderInput("bins",
                  "Number of bins:",
                  min = 1,
                  max = 50,
                  value = 30)
    ),

    # Show a plot of the generated distribution
    mainPanel(
      plotOutput("distPlot")
    )
  )
))
```

177

# server.R : Specifies what R is doing

```
library(shiny)

# Define server logic required to draw a histogram
shinyServer(function(input, output) {

  # Expression that generates a histogram. The expression is
  # wrapped in a call to renderPlot to indicate that:
  #

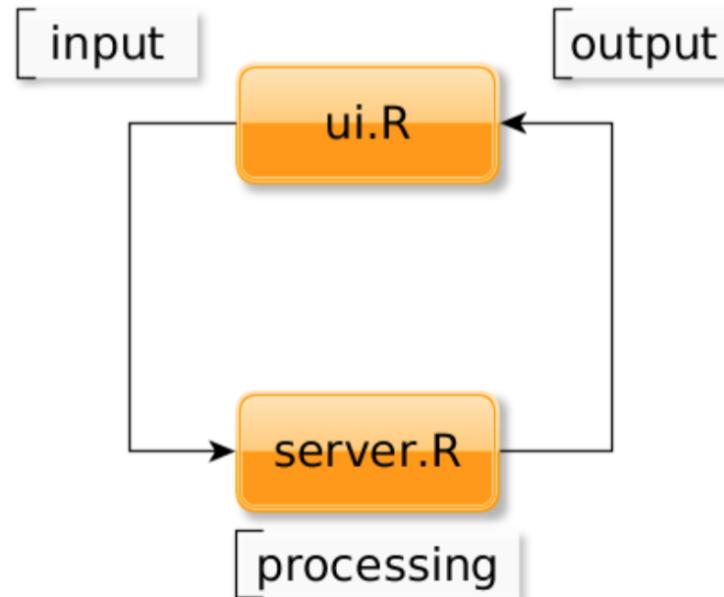
  # 1) It is "reactive" and therefore should be automatically
  #     re-executed when inputs change
  # 2) Its output type is a plot

  output$distPlot <- renderPlot({
    x      <- faithful[, 2] # Old Faithful Geyser data
    bins <- seq(min(x), max(x), length.out = input$bins + 1)

    # draw the histogram with the specified number of bins
    hist(x, breaks = bins, col = 'darkgray', border = 'white')
  })
})
```

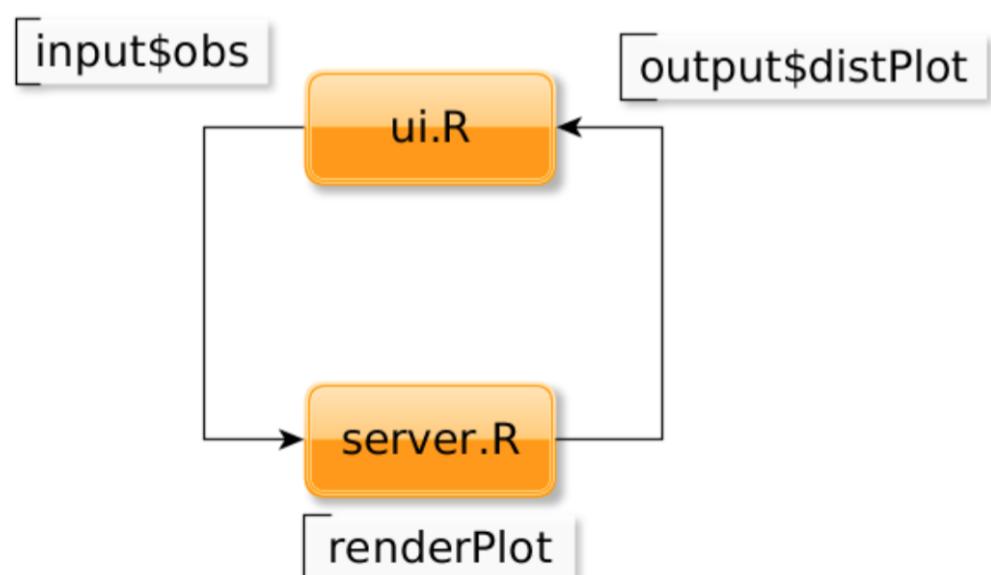
178

# Relationship of ui.R and server.R



179

# Relationship of ui.R and server.R



180

# For Spatial data

**sp, maptools** - Tools for loading and using spatial data including shapefiles.

**maps** - Easy to use map polygons for plots.

**ggmap** - Download street maps straight from Google maps and use them as a background in your ggplots.

181

# To write your own R packages

**devtools** - An essential suite of tools for turning your code into an R package.

**testthat** - testthat provides an easy way to write unit tests for your code projects.

**roxygen2** - A quick way to document your R packages. roxygen2 turns inline code comments into documentation pages and builds a package namespace.

182

# Hypothesis testing in R

## Fundamental of Hypothesis Testing

There two types of **statistical inferences**, **Estimation** and **Hypothesis Testing**

**Hypothesis Testing:** A hypothesis is a claim (assumption) about one or more population parameters.

- Average price of a six-pack in the U.S. is  $\mu = \$4.90$
- The population mean monthly cell phone bill of this city is:  $\mu = \$42$
- The average number of TV sets in U.S. Homes is equal to three;  $\mu = 3$



It Is always about a population parameter, not about a sample statistic

Sample evidence is used to assess the probability that the claim about the population parameter is true

**A. It starts with Null Hypothesis,  $H_0$**

$$H_0: \quad =$$

1. We begin with the assumption that  $H_0$  is true and any difference between the sample statistic and true population parameter is due to chance and not a real (systematic) difference.
2. Similar to the notion of “innocent until proven guilty”
3. That is, “innocence” is a null hypothesis.

185



Null Hypo, Continued

4. Refers to the status quo
5. Always contains “=” , “≤” or “≥” sign
6. May or may not be **rejected**

**B. Next we state the Alternative Hypothesis,  $H_1$**

1. Is the opposite of the null hypothesis
  1. e.g., The average number of TV sets in U.S. homes is not equal to 3 ( $H_1: \mu \neq 3$ )
2. Challenges the status quo
3. Never contains the “=” , “≤” or “≥” sign
4. May or may not be **proven**
5. Is generally the hypothesis that the researcher is trying to prove. Evidence is always examined with respect to  $H_1$ , never with respect to  $H_0$ .
6. We never “accept”  $H_0$ , we either “reject” or “not reject” it

186



## Summary:

- In the process of hypothesis testing, the null hypothesis initially is assumed to be true
- Data are gathered and examined to determine whether the evidence is strong enough with respect to the alternative hypothesis to reject the assumption.
- In other words, the burden is placed on the researcher to show, using sample information, that the null hypothesis is false.
- If the sample information is sufficient enough in favor of the alternative hypothesis, then the null hypothesis is rejected. This is the same as saying if the persecutor has enough evidence of guilt, the “innocence is rejected.”
- Of course, erroneous conclusions are possible, type I and type II errors.

187

## Reason for Rejecting $H_0$

Illustration: Let say, we assume that average age in the US is 50 years ( $H_0=50$ ). If in fact this is the true (unknown) population mean, it is unlikely that we get a sample mean of 20. So, if we have a sample that produces an average of 20, then we reject that the null hypothesis that average age is 50. (note that we are rejecting our assumption or claim). (would we get 20 if the true population mean was 50? NO. That is why we reject 50)

## How Is the Test done?

We use the distribution of a Test Statistic, such as Z or t as the criteria.

188



### A. Rejection Region Method:

Divide the distribution into rejection and non-rejection regions

Defines the unlikely values of the sample statistic if the null hypothesis is true, the critical value(s)

Defines **rejection region** of the sampling distribution

Rejection region(s) is designated by  $\alpha$ , (level of significance)

Typical values are .01, .05, or .10

$\alpha$  is selected by the researcher at the beginning

$\alpha$  provides the critical value(s) of the test

189

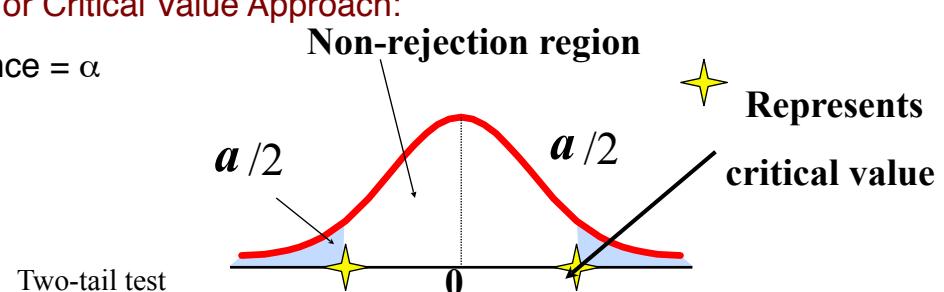


### Rejection Region or Critical Value Approach:

Level of significance =  $\alpha$

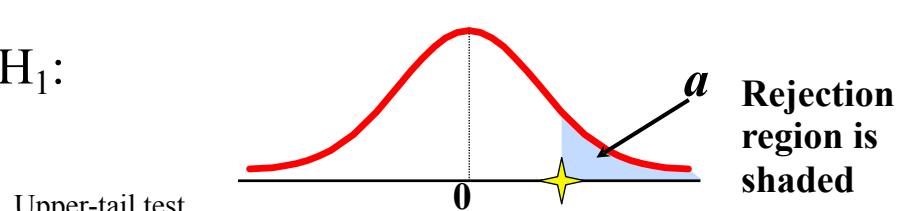
$$H_0: \mu = 12$$

$$H_1: \mu \neq 12$$



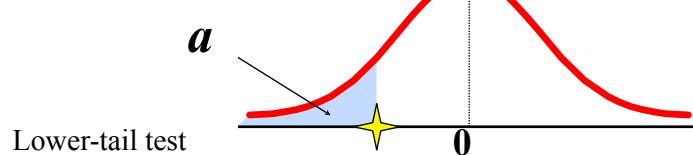
$$H_0: \mu \leq 12 \quad H_1:$$

$$\mu > 12$$



$$H_0: \mu \geq 12$$

$$H_1: \mu < 12$$



190



### P-Value Approach –

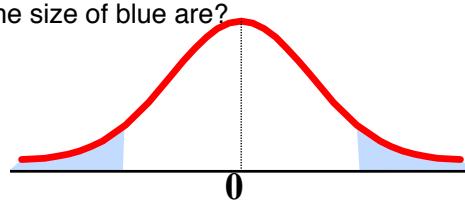
P-value=Max. Probability of (Type I Error), calculated from the sample.

Given the sample information what is the size of blue area?

$$H_0: \mu = 12$$

$$H_1: \mu \neq 12$$

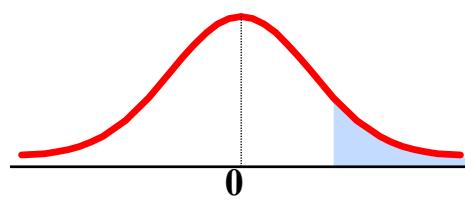
Two-tail test



$$H_0: \mu \leq 12 \quad H_1:$$

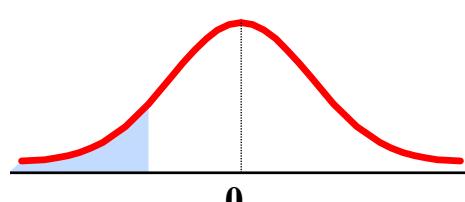
$$\mu > 12$$

Upper-tail test



$$H_0: \mu \geq 12$$

$$H_1: \mu < 12$$



191

## Type I and II Errors:

The size of  $\alpha$ , the rejection region, affects the risk of making different types of incorrect decisions.

### Type I Error

Rejecting a **true null hypothesis** when it should **NOT** be rejected

Considered a serious type of error

The probability of Type I Error is  $\alpha$

It is also called **level of significance** of the test

### Type II Error

Fail to reject a **false null hypothesis** that should have been rejected

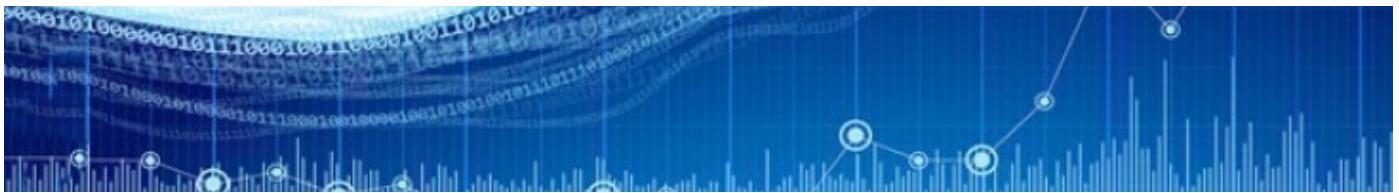
The probability of Type II Error is  $\beta$

192



Decision		Actual Situation		
		Hypothesis Testing	Legal System	
	H <sub>0</sub> True	H <sub>0</sub> False	Innocence	Not innocence
Do Not Reject H <sub>0</sub>	No Error (1 - α)	Type II Error (β)	No Error (not guilty, found not guilty) (1 - α)	Type II Error (guilty, found not guilty) (β)
Reject H <sub>0</sub>	Type I Error (α)	No Error (1 - β)	Type I Error (Not guilty, found guilty) (α)	No Error (guilty, found guilty) (1 - β)

193



Type I and Type II errors cannot happen at the same time

1. Type I error can only occur if H<sub>0</sub> is **true**
2. Type II error can only occur if H<sub>0</sub> is **false**
3. There is a tradeoff between type I and II errors. If the probability of type I error (  $\alpha$  ) increased, then the probability of type II error (  $\beta$  ) declines.
4. When the difference between the hypothesized parameter and the actual true value is small, the probability of type two error (the non-rejection region) is larger.
5. Increasing the sample size, n, for a given level of  $\alpha$ , reduces  $\beta$

194



## B. P-Value approach to Hypothesis Testing:

1. The rejection region approach allows you to examine evidence but restrict you to not more than a certain probability (say  $\alpha = 5\%$ ) of rejecting a true  $H_0$  by mistake.
2. The P-value approach allows you to use the information from the sample and then calculate the **maximum probability of rejecting a true  $H_0$  by mistake.**
3. Another way of looking at P-value is the probability of observing a sample information of “A=11.5” when the true population parameter is “12=B”. The P-value is the **maximum probability** of such mistake taking place.

195



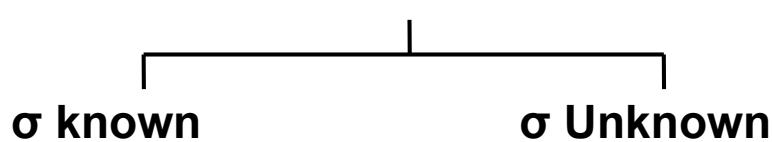
4. That is to say that P-value is the smallest value of  $\alpha$  for which  $H_0$  can be rejected based on the sample information
5. Convert Sample Statistic (e.g., sample mean) to Test Statistic (e.g., Z statistic )
6. Obtain the **p-value** from a table or computer
7. Compare the **p-value** with  $\alpha$

If  $p\text{-value} < \alpha$ , reject  $H_0$

If  $p\text{-value} \geq \alpha$ , do not reject  $H_0$

196

# Test of Hypothesis for the Mean



The test statistic is:

$$Z = \frac{\bar{X} - \mu}{\frac{\sigma}{\sqrt{n}}}$$

The test statistic is:

$$t_{n-1} = \frac{\bar{X} - \mu}{\frac{s}{\sqrt{n}}}$$

197

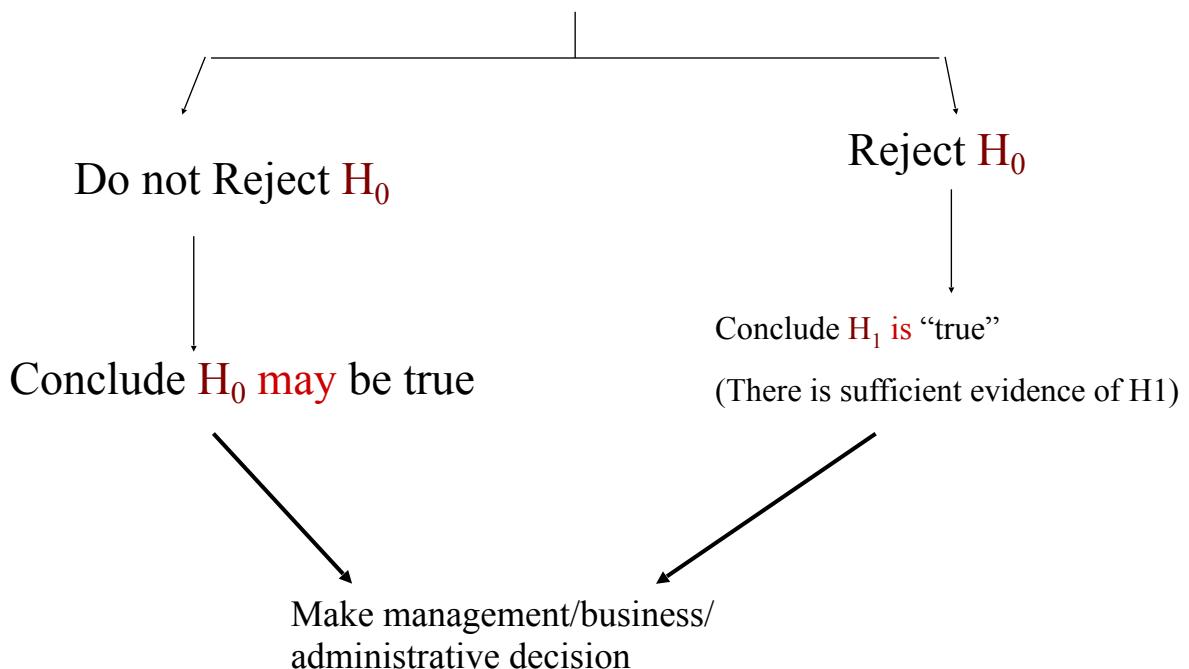
# Steps to Hypothesis Testing

1. State the  $H_0$  and  $H_1$  clearly
2. Identify the test statistic (two-tail, one-tail, and Z or t distribution)
3. Depending on the type of risk you are willing to take, specify the level of significance,
4. Find the decision rule, critical values, and rejection regions. If –  
 $CV < \text{actual value (sample statistic)} < +CV$ , then **do not reject the  $H_0$**
5. Collect the data and do the calculation for the actual values of the test statistic from the sample

198



Make statistical decision



199



## When do we use a two-tail test?

The answer depends on the question you are trying to answer.

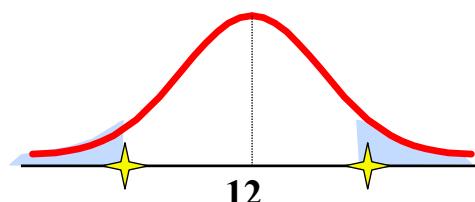
A two-tail is used when the researcher has no idea which direction the study will go, interested in both direction. (example: testing a new technique, a new product, a new theory and we don't know the direction)

A new machine is producing 12 fluid once can of soft drink. The quality control manager is concern with cans containing too much or too little.

Then, the test is a two-tailed test. That is the two rejection regions in tails is most likely (higher probability) to provide evidence of H<sub>1</sub>.

$$H_0 : \mu = 12 \text{ oz}$$

$$H_1 : \mu \neq 12 \text{ oz}$$



200



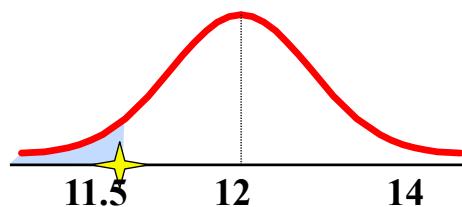
One-tail test is used when the researcher is interested in the direction.

Example: The soft-drink company puts a label on cans claiming they contain 12 oz. A consumer advocate desires to test this statement. She would assume that each can contains **at least** 12 oz and tries to find evidence to the contrary. That is, she examines the evidence for less than 12 oz.

What tail of the distribution is the most logical (higher probability) to find that evidence? The only way to reject the claim is to get evidence of less than 12 oz, left tail.

$$H_0 : \mu \geq 12 \text{ oz}$$

$$H_1 : \mu < 12 \text{ oz}$$



201

## Review of Hypo. Testing

What is HT?

Probability of making erroneous conclusions

Type I – only when Null Hypo is true

Type II – only when Null Hypo is false

Two Approaches

The Rejection or Critical Value Approach

The P-value Approach (we calculate the observed level of significance)

Test Statistics

Z- distribution if Population Std. Dev. is Known

t-distribution if the Population Std. Dev. is unknown

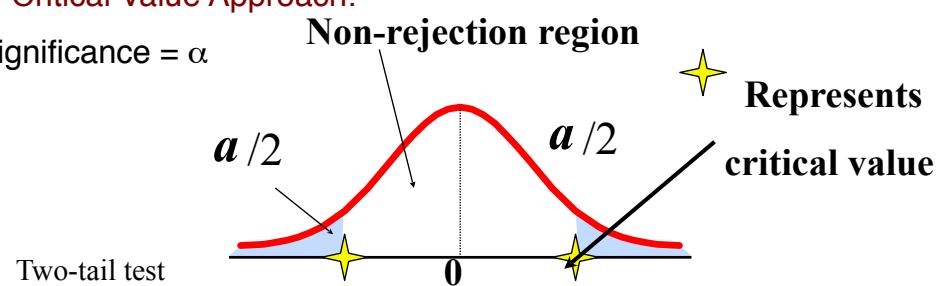
202

### Rejection Region or Critical Value Approach:

The given level of significance =  $\alpha$

$$H_0: \mu = 12$$

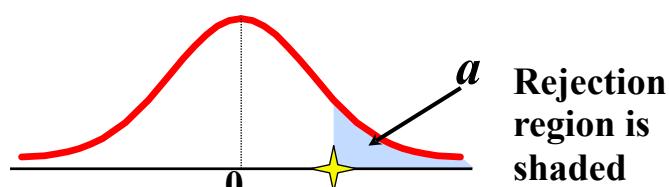
$$H_1: \mu \neq 12$$



$$H_0: \mu \leq 12 \quad H_1:$$

$$\mu > 12$$

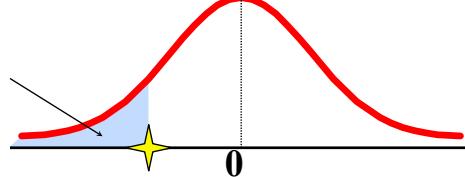
Upper-tail test



$$H_0: \mu \geq 12$$

$$H_1: \mu < 12$$

Lower-tail test



203

### P-Value Approach –

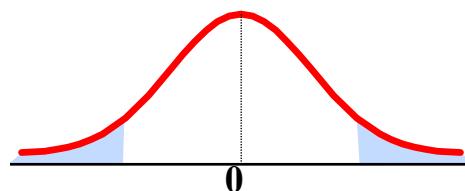
P-value=Max. Probability of (Type I Error), calculated from the sample.

Given the sample information what is the size of the blue areas? (The observed level of significance)

$$H_0: \mu = 12$$

$$H_1: \mu \neq 12$$

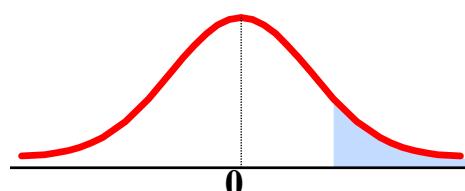
Two-tail test



$$H_0: \mu \leq 12 \quad H_1:$$

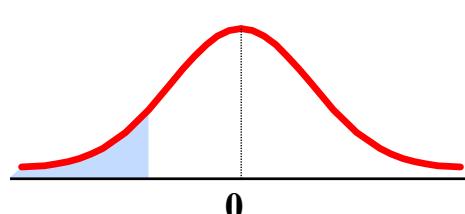
$$\mu > 12$$

Upper-tail test



$$H_0: \mu \geq 12$$

$$H_1: \mu < 12$$



204



Example 1:

Let's assume a sample of 25 cans produced a sample mean of 11.5 0z and the population std dev=1 0z.

**Question 1:**

At a 5% level of significance (that is allowing for a maximum of 5% prob. of rejecting a true null hypo), is there evidence that the population mean is different from 12 oz?

Null Hypo is:?

Alternative Hypo is?

Can both approaches be used to answer this question?

205



**A: Rejection region approach:** calculate the actual test statistics and compare it with the critical values

**B: P-value approach:** calculate the actual probability of type I error given the sample information. Then compare it with 1%, 5%, or 10% level of significance.

Interpretation of Critical Value/Rejection Region Approach:

Interpretation of P-value Approach:

206



### Question 2:

At a 5% level of significance (that is allowing for a maximum of 5% prob. of rejecting a true null hypo), is the evidence that the population mean is **less than 12 oz**?

Null Hypo is:?

Alternative Hypo is?

Can both approaches be used to answer this question?

Interpretation of Critical Value Approach:

Interpretation of P-value Approach:

207



### Question 3:

If in fact the pop. mean is 12 oz, what is the probability of obtaining a sample mean of 11.5 or less oz (sample size 25)? Null

Null Hypo is:?

Alternative Hypo is?

### Question 4:

If in fact the pop. mean is 12 oz, and the sample mean is 11.5 (or less), what is the probability of erroneously rejecting the null hypo that the pop. mean is 12 oz?

Null Hypo is:?

Alternative Hypo is?

Can both approaches be used to answer these question?

208

# Connection to Confidence Intervals

While the confidence interval estimation and hypothesis testing serve different purposes, they are based on same concept and conclusions reached by two methods are consistent for a two-tail test.

In CI method we estimate an interval for the population mean with a degree of confidence. If the estimated interval **contains** the hypothesized value under the hypothesis testing, then this is equivalent of **not rejecting** the null hypothesis. For example: for the beer sample with mean 5.20, the confidence interval is:

$$P(4.61 \leq \mu \leq 5.78) = 95\%$$

Since this interval contains the Hypothesized mean (\$4.90), we do not (did not) reject the null hypothesis at  $\alpha = .05$

Did not reject and within the interval, thus consistent results.

209

# Introduction to Regression Analysis

**Regression analysis** is used to:

Predict values of a dependent variable, Y, based on its relationship with values of at least one independent variable, X.

Explain the impact of changes in an independent variable on the dependent variable by estimating the **numerical value** of the relationship

**Dependent variable:** the variable we wish to explain

**Independent variable:** the variable used to explain the dependent variable

210

# Simple Linear Regression Model

Only **one** independent variable (thus, simple), X

Relationship between X and Y is described by a linear function

Changes in Y are assumed to be caused by changes in X, that is,

Change In X → **Causes** → Change in Y

211

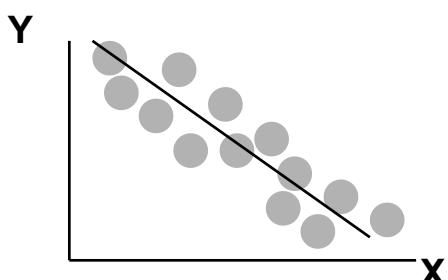
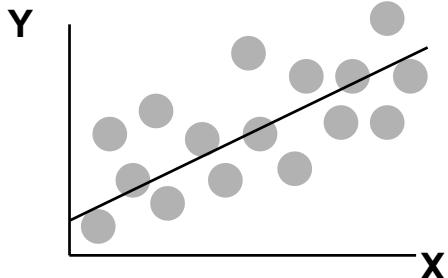
## Important points before we start a regression analysis:

- The most important thing in deciding whether or not there is a relationship between X and Y is to have a systematic model that is based on logical reasons.
- Investigate the nature of the relationship between X and Y (use scatter diagram, covariance, correlation coefficient)
- Remember that regression is not an exact or deterministic mathematical equation. It is a **behavioral relationship** that is subject to randomness.
- Remember that X is not the only thing that explains the behavior of Y. There are other factor that you may not have information about.
- All you are trying to do is to have an estimate of the relationship using the **best linear fit** possible

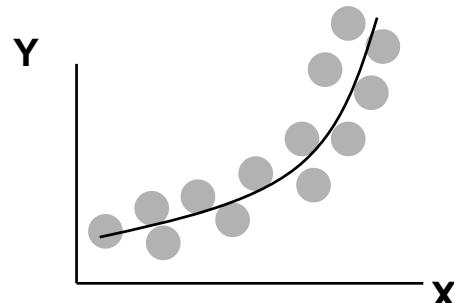
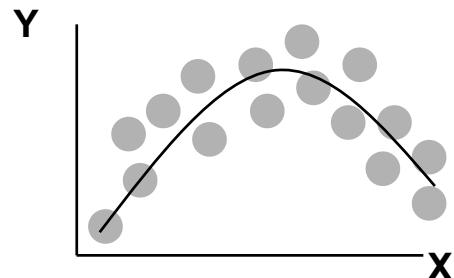
212

# Types of Relationships

## Linear relationships



## Curvilinear relationships

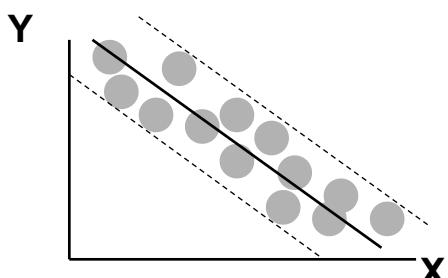
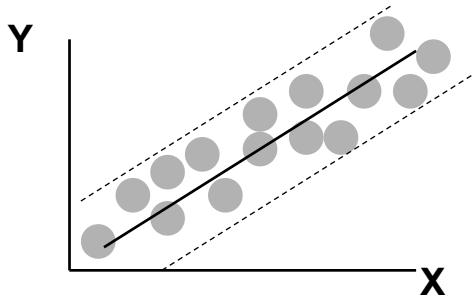


213

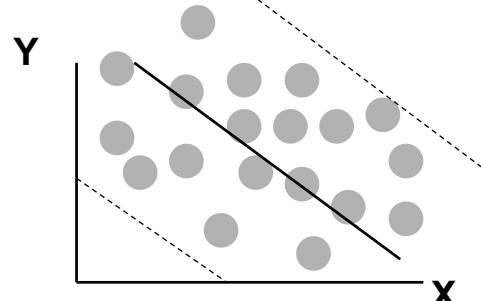
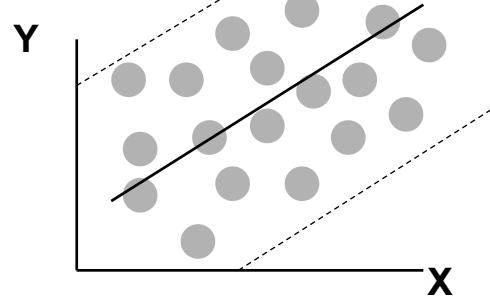
# Types of Relationships

(continued)

## Strong relationships



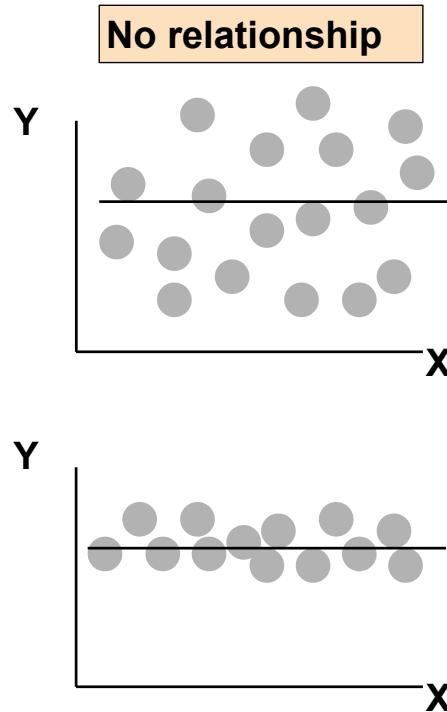
## Weak relationships



214

# Types of Relationships

(continued)



215

## Simple Linear Regression Conceptual Model

The population regression model: This is a conceptual model, a hypothesis, or a postulation

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

Annotations for the equation:

- Dependent Variable: Points to the  $Y_i$  term.
- Population Y intercept: Points to the  $\beta_0$  term.
- Population Slope Coefficient: Points to the  $\beta_1$  term.
- Independent Variable: Points to the  $X_i$  term.
- Random Error term: Points to the  $\epsilon_i$  term.
- Linear component: Brackets under the terms  $\beta_0 + \beta_1 X_i$ .
- Random Error component: Brackets under the term  $\epsilon_i$ .

216



- The model to be estimated from sample data is:

$$Y_i = b_0 + b_1 X_i + e_i$$

Residual  
(random  
error from  
the sample)

- The actual estimated from the sample

Estimated (or predicted) Y value for observation i

Estimate of the regression intercept

Estimate of the regression slope

Value of X for observation i

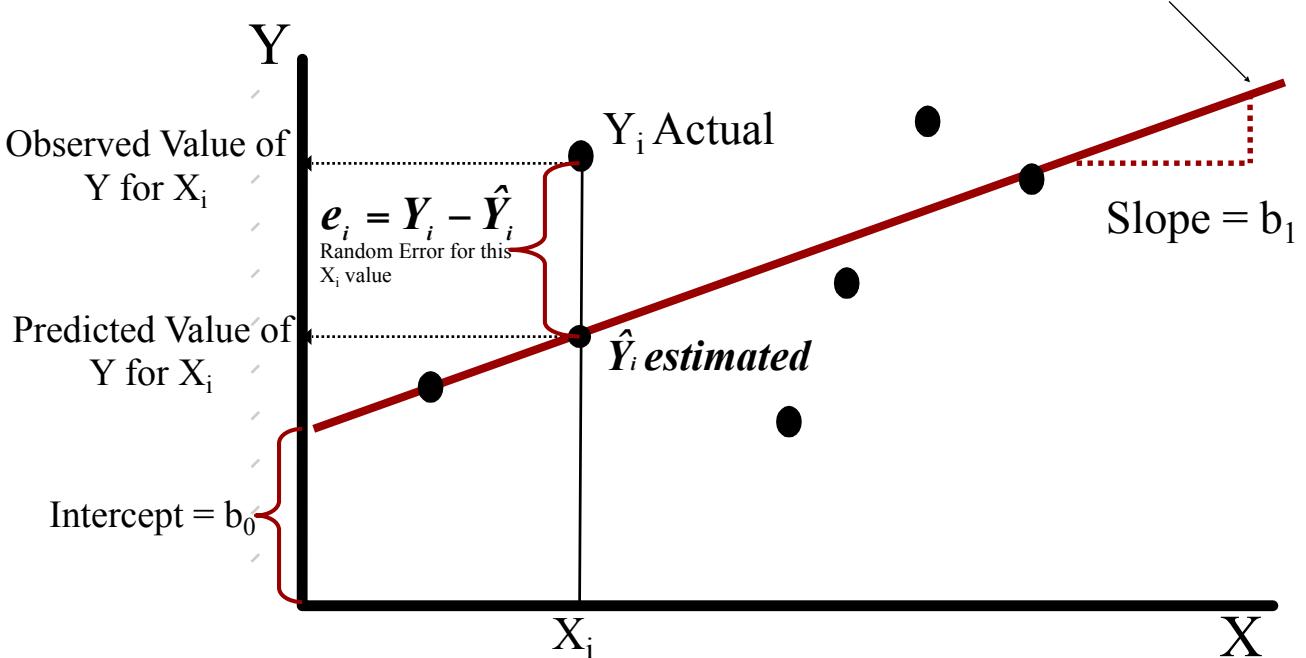
$$\hat{Y}_i = b_0 + b_1 X_i$$

– Where  $e_i = Y_i - \hat{Y}_i$

217

## Simple Linear Regression Model

$$\hat{Y}_i = b_0 + b_1 X_i$$



218

# Interpretation of the slope and the intercept

$$\beta_0 = E(Y | X = 0) ; \quad \beta_1 = \Delta E(Y|X)/\Delta (X);$$

$\beta_0$  is the estimated average value of Y when the value of X is  $b_0$  zero

$\beta_1$  is the estimated change in the average value of Y as a result of a one-unit change in X

Units of measurement of X and Y are very important for the correct interpretation of the slope and the intercept

Example:  $\widehat{App\ Val} = 165.03 + 6.93 (Lot\ size)$

Predict the app. Value of a house with 10,000 s.f. lot size

$$\widehat{App\ Val} = 165.03 + 6.93 (10) = \$234,330$$

How Good is this prediction?

219

# How Good is the Model's prediction Power?

Total variation is made up of two parts:

$$SST = SSR + SSE$$

Total Sum of Squares

Regression Sum of Squares

Error Sum of Squares

$$SST = \sum (Y_i - \bar{Y})^2 \quad SSR = \sum (\hat{Y}_i - \bar{Y})^2 \quad SSE = \sum (Y_i - \hat{Y}_i)^2$$

where:

$\bar{Y}$  = Average value of the dependent variable

$Y_i$  = Observed values of the dependent variable

$\hat{Y}_i$  = Predicted value of Y for the given  $X_i$  value

220



SST = total sum of squares

Measures total variation of the  $Y_i$  values around their mean

SSR = regression sum of squares (Explained)

Explained portion of total variation attributed to Y's

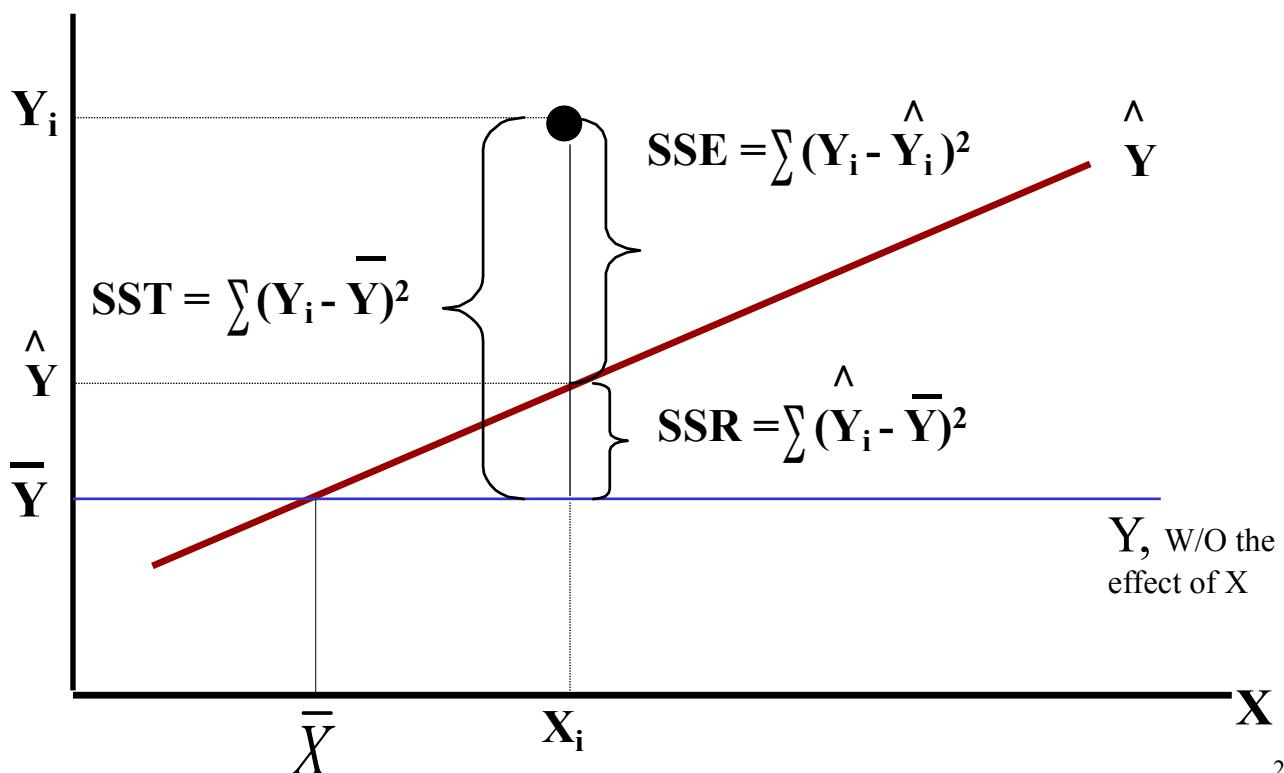
relationship with X

SSE = error sum of squares (Unexplained)

Variation of Y values attributable to other factors than its

relationship with X

221



222

# How Good is the Model's prediction Power?

The coefficient of determination is the portion of the total variation in the dependent variable, Y, that is explained by variation in the independent variable, X

The coefficient of determination is also called r-squared and is denoted as  $r^2$

$$r^2 = \frac{SSR}{SST} = \frac{\text{regression sum of squares}}{\text{total sum of squares}}$$

$$0 \leq r^2 \leq 1$$

223

## Standard Error of Estimate

The standard deviation of the variation of observations around the regression line is estimated by

Where SSE = error sum of squares; n = sample size

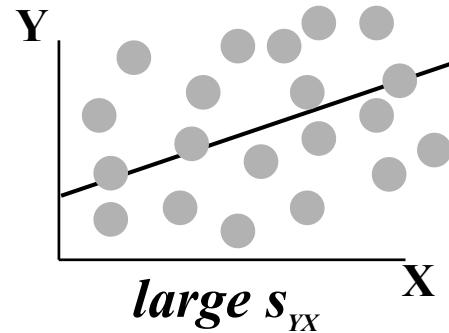
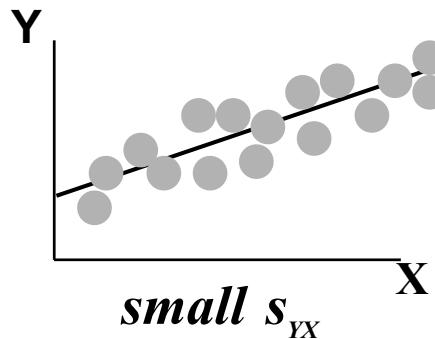
$$S_{yx} = \sqrt{\frac{SSE}{n-2}} = \sqrt{\frac{\sum_{i=1}^n (Y_i - \hat{Y}_i)^2}{n-2}} = \sqrt{MSE}$$

The concept is the same as the standard deviation (average difference) around the mean of a univariate

224

# Comparing Standard Errors

$S_{YX}$  is a measure of the variation of observed Y values from the regression line



The magnitude of  $S_{YX}$  should always be judged relative to the size of the Y values in the sample data

i.e.,  $S_{YX} = \$36.34K$  is moderately small relative to house prices in the \$200 - \$300K range (average 215K)

225

# Assumptions of Regression

## Normality of Error

Error values ( $\varepsilon$ ) are normally distributed for any given value of X

## Homoscedasticity

The probability distribution of the errors has constant variance

## Independence of Errors

Error values are statistically independent

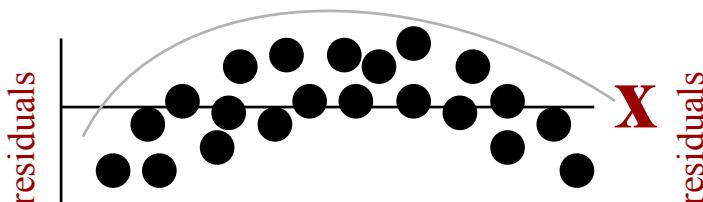
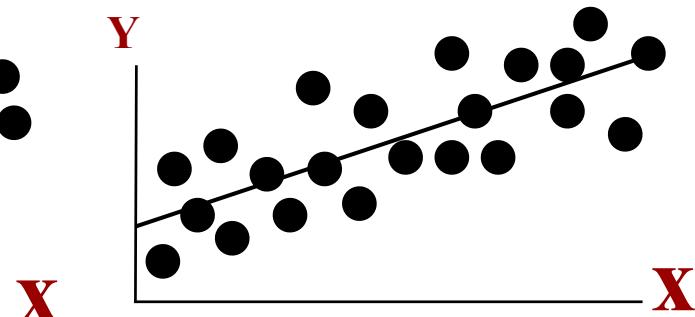
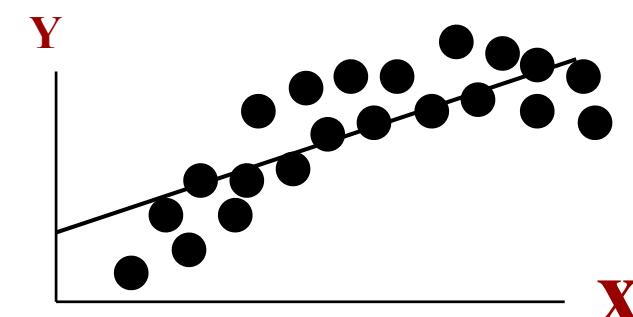
226

## How to investigate the appropriateness of the fitted model

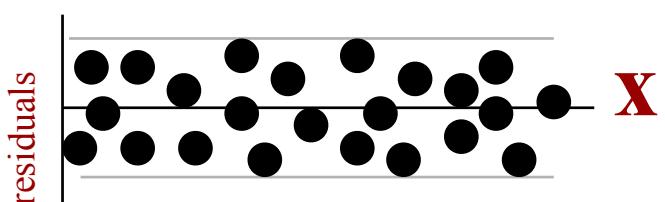
- The residual for observation  $i$ ,  $e_i$ , is the difference between its observed and predicted value;
- Check the assumptions of regression by examining the residuals  $e_i = Y_i - \hat{Y}_i$ 
  - Examine for linearity assumption
  - Examine for constant variance for all levels of  $X$  (homoscedasticity)
  - Evaluate normal distribution assumption
  - Evaluate independence assumption
- Graphical Analysis of Residuals
  - Can plot residuals vs.  $X$

227

## Residual Analysis for Linearity



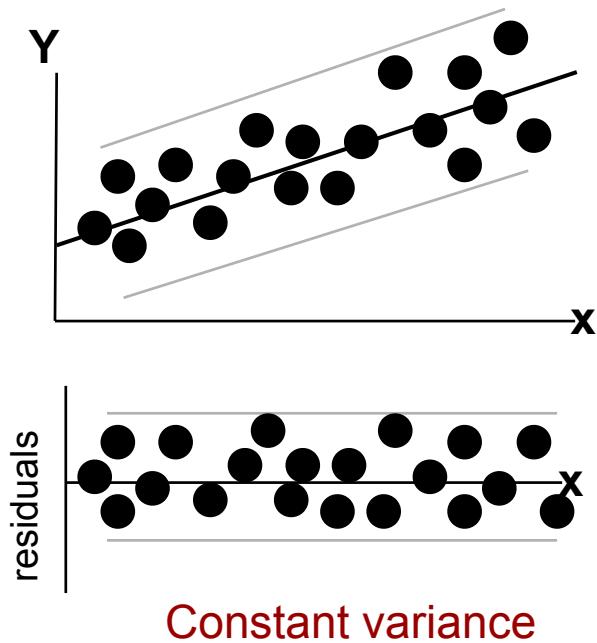
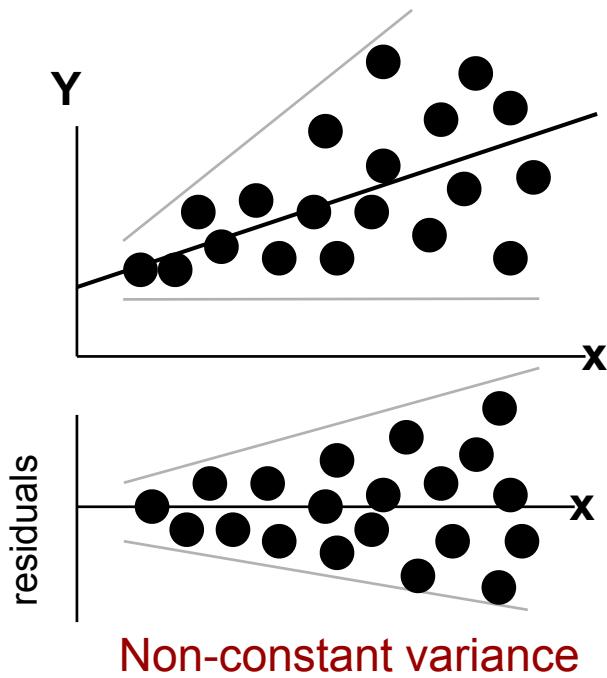
**Not Linear**



**Linear**

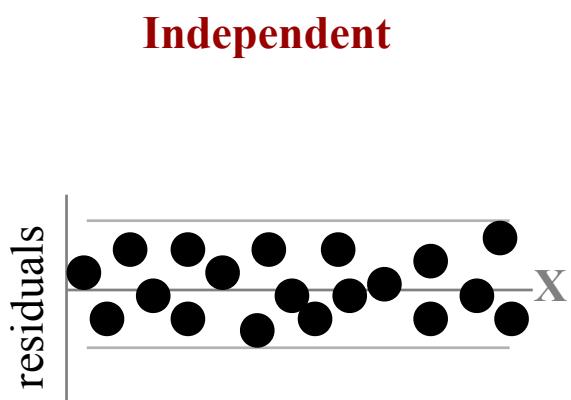
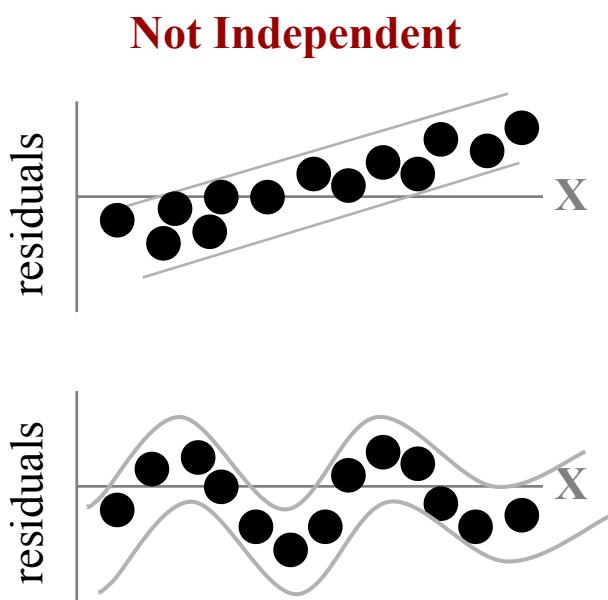
228

# Residual Analysis for Homoscedasticity



229

# Residual Analysis for Independence



230

# Testing differences in mean between two groups

Let's begin by loading the packages we'll need to get started

```
library(MASS)
library(plyr)
library(ggplot2)
```

```
# Rename the columns to have more descriptive names
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",
  "race", "mother.smokes", "previous.prem.labor", "hypertension", "uterine.irr",
  "physician.visits", "birthwt.grams")

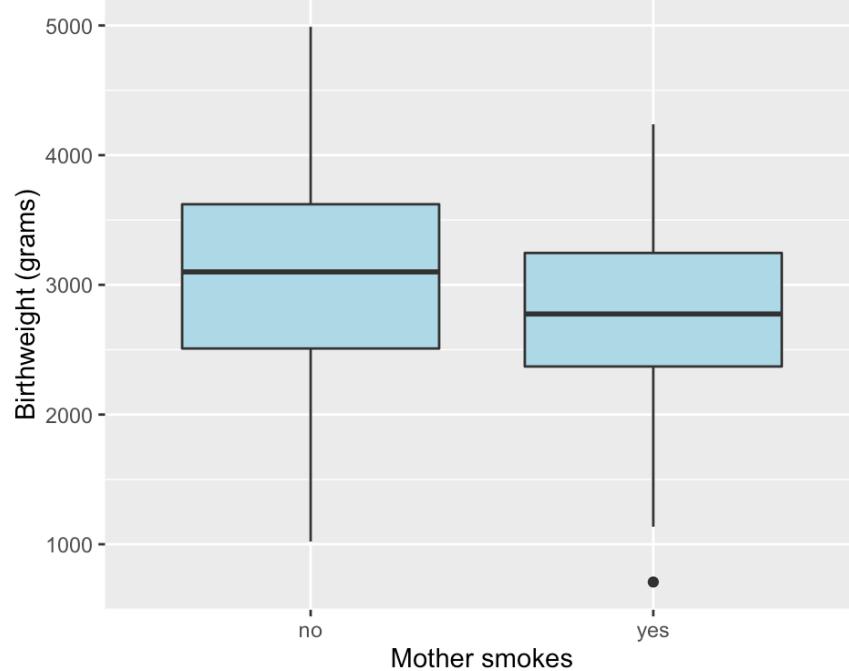
# Transform variables to factors with descriptive levels
birthwt <- transform(birthwt,
  race = as.factor(mapvalues(race, c(1, 2, 3),
    c("white", "black", "other"))),
  mother.smokes = as.factor(mapvalues(mother.smokes,
    c(0,1), c("no", "yes"))),
  hypertension = as.factor(mapvalues(hypertension,
    c(0,1), c("no", "yes"))),
  uterine.irr = as.factor(mapvalues(uterine.irr,
    c(0,1), c("no", "yes"))))
)
```

231

To start, it always helps to plot things

```
# Create boxplot showing how birthwt.grams varies between
# smoking status
qplot(x = mother.smokes, y = birthwt.grams,
  geom = "boxplot", data = birthwt,
  xlab = "Mother smokes",
  ylab = "Birthweight (grams)",
  fill = I("lightblue"))
```

232



This plot suggests that smoking is associated with lower birth weight.

233

## How can we assess whether this difference is statistically significant?

Let's compute a summary table

```
aggregate(birthwt.grams ~ mother.smokes, data = birthwt,
          FUN = function(x) {c(mean = mean(x), sd = sd(x))})
```

	mother.smokes	birthwt.grams.mean	birthwt.grams.sd
## 1	no	3055.6957	752.6566
## 2	yes	2771.9189	659.6349

234



The standard deviation is good to have, but to assess statistical significance we really want to have the standard error (which is the standard deviation adjusted by the group size).

```
aggregate(birthwt.grams ~ mother.smokes, data = birthwt,
          FUN = function(x) {c(mean = mean(x),
                                 se = sd(x) / sqrt(length(x))))})
```

```
##   mother.smokes birthwt.grams.mean birthwt.grams.se
## 1       no           3055.69565     70.18559
## 2      yes          2771.91892     76.68100
```

235

## t-test via t.test()

This difference is looking quite significant. To run a two-sample t-test, we can simply use the `t.test()` function.

```
birthwt.t.test <- t.test(birthwt.grams ~ mother.smokes, data = birthwt)
birthwt.t.test
```

```
##
##  Welch Two Sample t-test
##
## data: birthwt.grams by mother.smokes
## t = 2.7299, df = 170.1, p-value = 0.007003
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##    78.57486 488.97860
## sample estimates:
## mean in group no mean in group yes
##            3055.696           2771.919
```

236

# p-value

```
birthwt.t.test$p.value    # p-value  
  
## [1] 0.007002548  
  
birthwt.t.test$estimate   # group means  
  
## mean in group no mean in group yes  
##          3055.696           2771.919  
  
birthwt.t.test$conf.int   # confidence interval for difference  
  
## [1] 78.57486 488.97860  
## attr(,"conf.level")  
## [1] 0.95
```

237

## Workshop 8: t-test

Testing means between two groups

(a) Using the Cars93 data and the t.test() function, run a t-test to see if average MPG.highway is different between US and non-US vehicles.

Try doing this both using the formula style input and the x, y style input.

(b) What is the confidence interval for the difference?

(c) Repeat part (a) using the wilcox.test() function.

(d) Are your results for (a) and (c) very different?

238



# ANOVA and Regression



## Prepare data

```
library(MASS)
library(plyr)
library(ggplot2)
library(knitr)

# Rename the columns to have more descriptive names
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",
"race", "mother.smokes", "previous.prem.labor", "hypertension", "uterine.irr",
"physician.visits", "birthwt.grams")

# Transform variables to factors with descriptive levels
birthwt <- transform(birthwt,
  race = as.factor(mapvalues(race, c(1, 2, 3),
    c("white", "black", "other"))),
  mother.smokes = as.factor(mapvalues(mother.smokes,
    c(0,1), c("no", "yes"))),
  hypertension = as.factor(mapvalues(hypertension,
    c(0,1), c("no", "yes"))),
  uterine.irr = as.factor(mapvalues(uterine.irr,
    c(0,1), c("no", "yes")))
)
```

# One-way ANOVA example

Question: Is there a significant association between race and birthweight?

Here's a table showing the mean and standard error of birthweight by race.

```
aggregate(birthwt.grams ~ race, data = birthwt, FUN = mean)

##      race birthwt.grams
## 1 black     2719.692
## 2 other     2805.284
## 3 white    3102.719
```

241

**Terminology:** a  $k$ -way ANOVA is used to assess whether the mean of an outcome variable is constant across all combinations of  $k$  factors. The most common examples are 1-way ANOVA (looking at a single factor), and 2-way ANOVA (looking at two factors).

We'll use the `aov()` function. For convenience, `aov()` allows you to specify a formula.

```
summary(aov(birthwt.grams ~ race, data = birthwt))

##              Df  Sum Sq Mean Sq F value    Pr(>F)
## race          2  5015725 2507863   4.913 0.00834 ***
## Residuals   186  94953931  510505
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

242

# Workshop 9: ANOVA

Let's form our favourite birthwt data set.

```
# Rename the columns to have more descriptive names
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",
  "race", "mother.smokes", "previous.prem.labor", "hypertension", "uterine.irr",
  "physician.visits", "birthwt.grams")

# Transform variables to factors with descriptive levels
birthwt <- transform(birthwt,
  race = as.factor(mapvalues(race, c(1, 2, 3),
    c("white", "black", "other"))),
  mother.smokes = as.factor(mapvalues(mother.smokes,
    c(0,1), c("no", "yes"))),
  hypertension = as.factor(mapvalues(hypertension,
    c(0,1), c("no", "yes"))),
  uterine.irr = as.factor(mapvalues(uterine.irr,
    c(0,1), c("no", "yes"))))
)
```

243

- (a) Create a new factor that categorizes the number of physician visits into three levels: 0, 1, 2, 3 or more.

Hint: One way of doing this is with mapvalues, by mapping all instances of 3, 4,... etc, to “3 or more”.

- (b) Run an ANOVA to determine whether the average birth weight varies across number of physician visits.

244

# Linear Regression

Let's first import the data set and see what we're working with.

```
# Import data set
crime <- read.table("crime_simple.txt", sep = "\t", header = TRUE)
```

The variable names that this data set comes with are very confusing, and even misleading.

R: Crime rate: # of offenses reported to police per million population

Age: The number of males of age 14-24 per 1000 population

S: Indicator variable for Southern states (0 = No, 1 = Yes)

Ed: Mean # of years of schooling x 10 for persons of age 25 or older

Ex0: 1960 per capita expenditure on police by state and local government

Ex1: 1959 per capita expenditure on police by state and local government

LF: Labor force participation rate per 1000 civilian urban males age 14-24

M: The number of males per 1000 females

N: State population size in hundred thousands

NW: The number of non-whites per 1000 population

U1: Unemployment rate of urban males per 1000 of age 14-24

U2: Unemployment rate of urban males per 1000 of age 35-39

W: Median value of transferable goods and assets or family income in tens of \$

X: The number of families per 1000 earning below 1/2 the median income

245



We really need to give these variables better names

```
# Assign more meaningful variable names
colnames(crime) <- c("crime.per.million", "young.males", "is.south", "average.ed",
                      "exp.per.cap.1960", "exp.per.cap.1959", "labour.part",
                      "male.per.fem", "population", "nonwhite",
                      "unemp.youth", "unemp.adult", "median.assets", "num.low.salary")

# Convert is.south to a factor
# Divide average.ed by 10 so that the variable is actually average education
# Convert median assets to 1000's of dollars instead of 10's
crime <- transform(crime, is.south = as.factor(is.south),
                    average.ed = average.ed / 10,
                    median.assets = median.assets / 100)

# print summary of the data
summary(crime)
```

246

```

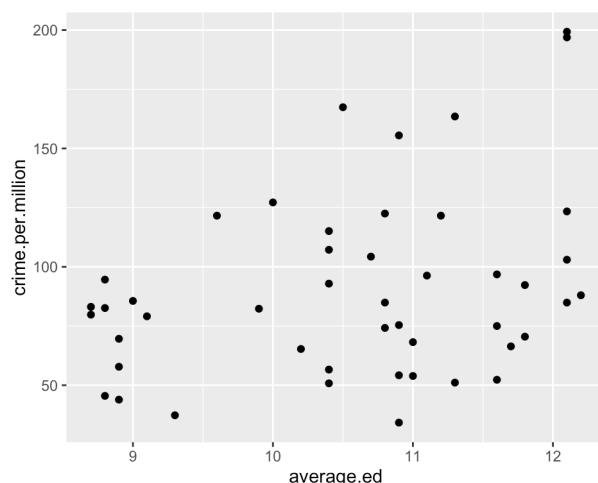
##   crime.per.million   young.males    is.south average.ed
##   Min. : 34.20      Min. :119.0     0:31      Min. : 8.70
##   1st Qu.: 65.85     1st Qu.:130.0    1:16      1st Qu.: 9.75
##   Median : 83.10     Median :136.0    Median :10.80
##   Mean   : 90.51     Mean   :138.6     Mean   :10.56
##   3rd Qu.:105.75     3rd Qu.:146.0    3rd Qu.:11.45
##   Max.  :199.30     Max.  :177.0     Max.  :12.20
##   exp.per.cap.1960  exp.per.cap.1959  labour.part male.per.fem
##   Min. : 45.0       Min. : 41.00    Min. :480.0   Min. : 934.0
##   1st Qu.: 62.5      1st Qu.: 58.50   1st Qu.:530.5  1st Qu.: 964.5
##   Median : 78.0      Median : 73.00   Median :560.0   Median : 977.0
##   Mean   : 85.0      Mean   : 80.23   Mean   :561.2   Mean   : 983.0
##   3rd Qu.:104.5      3rd Qu.: 97.00   3rd Qu.:593.0  3rd Qu.: 992.0
##   Max.  :166.0       Max.  :157.00   Max.  :641.0   Max.  :1071.0
##   population        nonwhite      unemp.youth unemp.adult
##   Min. : 3.00       Min. : 2.0     Min. : 70.00  Min. :20.00
##   1st Qu.: 10.00     1st Qu.: 24.0   1st Qu.: 80.50 1st Qu.:27.50
##   Median : 25.00     Median : 76.0   Median : 92.00  Median :34.00
##   Mean   : 36.62     Mean   :101.1   Mean   : 95.47  Mean   :33.98
##   3rd Qu.: 41.50     3rd Qu.:132.5   3rd Qu.:104.00 3rd Qu.:38.50
##   Max.  :168.00     Max.  :423.0   Max.  :142.00  Max.  :58.00
##   median.assets    num.low.salary
##   Min. :2.880       Min. :126.0
##   1st Qu.:4.595      1st Qu.:165.5
##   Median :5.370      Median :176.0
##   Mean   :5.254      Mean   :194.0
##   3rd Qu.:5.915      3rd Qu.:227.5
##   Max.  :6.890       Max.  :276.0

```

247

## First step: some plotting and summary statistics

```
# Scatter plot of outcome (crime.per.million) against average.ed
qplot(average.ed, crime.per.million, data = crime)
```

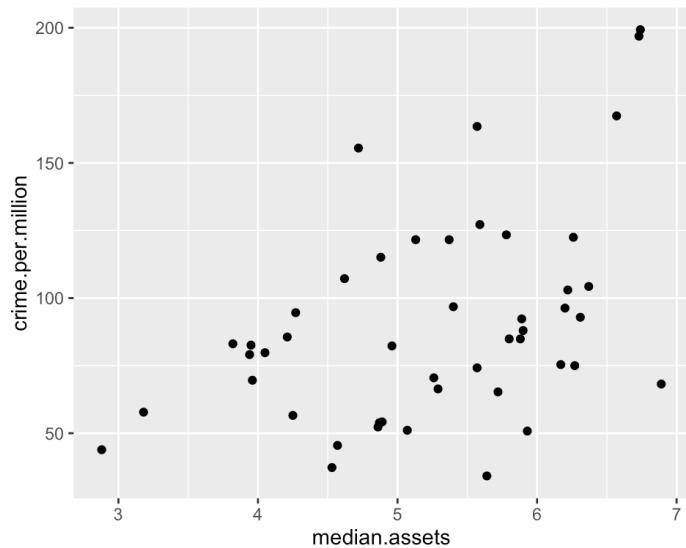


```
# correlation between education and crime
with(crime, cor(average.ed, crime.per.million))
```

```
## [1] 0.3228349
```

248

```
# Scatter plot of outcome (crime.per.million) against median.assets  
qplot(median.assets, crime.per.million, data = crime)
```

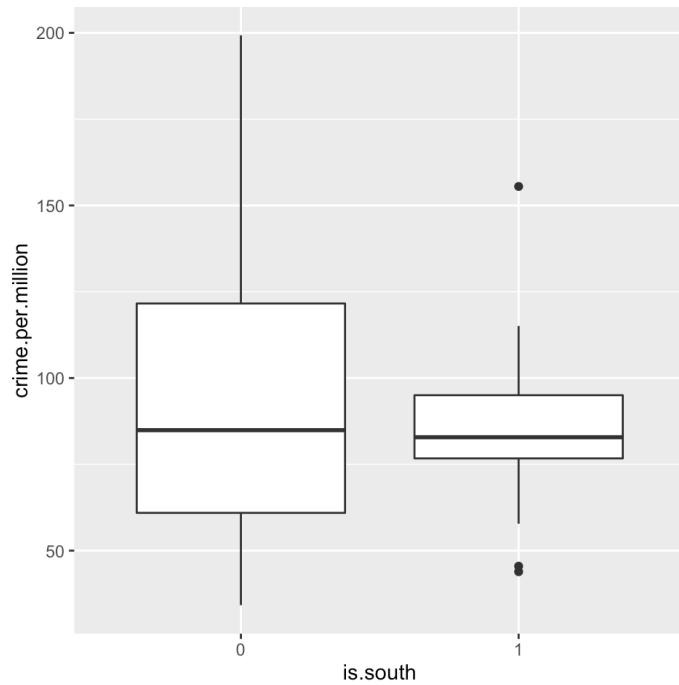


```
# correlation between education and crime  
with(crime, cor(median.assets, crime.per.million))
```

```
## [1] 0.4413199
```

249

```
# Boxplots showing crime rate broken down by southern vs non-southern state  
qplot(is.south, crime.per.million, geom = "boxplot", data = crime)
```



250

# Constructing a regression model

```
crime.lm <- lm(crime.per.million ~ ., data = crime)
# Summary of the linear regression model
crime.lm
```

```
##
## Call:
## lm(formula = crime.per.million ~ ., data = crime)
##
## Coefficients:
## (Intercept) young.males    is.south1   average.ed
## -6.918e+02   1.040e+00   -8.308e+00   1.802e+01
## exp.per.cap.1960 exp.per.cap.1959 labour.part male.per.fem
## 1.608e+00     -6.673e-01   -4.103e-02   1.648e-01
## population      nonwhite    unemp.youth unemp.adult
## -4.128e-02     7.175e-03   -6.017e-01   1.792e+00
## median.assets   num.low.salary
## 1.374e+01      7.929e-01
```

251

```
summary(crime.lm)
```

```
##
## Call:
## lm(formula = crime.per.million ~ ., data = crime)
##
## Residuals:
##   Min     1Q Median     3Q    Max
## -34.884 -11.923 -1.135 13.495 50.560
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t| )
## (Intercept) -6.918e+02 1.559e+02 -4.438 9.56e-05 ***
## young.males  1.040e+00 4.227e-01  2.460 0.01931 *
## is.south1    -8.308e+00 1.491e+01 -0.557 0.58117
## average.ed   1.802e+01 6.497e+00  2.773 0.00906 **
## exp.per.cap.1960 1.608e+00 1.059e+00  1.519 0.13836
## exp.per.cap.1959 -6.673e-01 1.149e+00 -0.581 0.56529
## labour.part   -4.103e-02 1.535e-01 -0.267 0.79087
## male.per.fem  1.648e-01 2.099e-01  0.785 0.43806
## population    -4.128e-02 1.295e-01 -0.319 0.75196
## nonwhite       7.175e-03 6.387e-02  0.112 0.91124
## unemp.youth   -6.017e-01 4.372e-01 -1.376 0.17798
## unemp.adult   1.792e+00 8.561e-01  2.093 0.04407 *
## median.assets 1.374e+01 1.058e+01  1.298 0.20332
## num.low.salary 7.929e-01 2.351e-01  3.373 0.00191 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.94 on 33 degrees of freedom
## Multiple R-squared:  0.7692, Adjusted R-squared:  0.6783
## F-statistic: 8.462 on 13 and 33 DF,  p-value: 3.686e-07
```

252

```

options(scipen=4) # Set scipen = 0 to get back to default

summary(crime.lm)

## 
## Call:
## lm(formula = crime.per.million ~ ., data = crime)
##
## Residuals:
##    Min     1Q Median     3Q    Max 
## -34.884 -11.923 -1.135 13.495 50.560 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -691.837588 155.887918 -4.438 0.0000956 ***
## young.males   1.039810  0.422708  2.460  0.01931 *  
## is.south1     -8.308313 14.911588 -0.557  0.58117    
## average.ed    18.016011  6.496504  2.773  0.00906 ** 
## exp.per.cap.1960 1.607818  1.058667  1.519  0.13836    
## exp.per.cap.1959 -0.667258 1.148773 -0.581  0.56529    
## labour.part   -0.041031  0.153477 -0.267  0.79087    
## male.per.fem   0.164795  0.209932  0.785  0.43806    
## population    -0.041277  0.129516 -0.319  0.75196    
## nonwhite       0.007175  0.063867  0.112  0.91124    
## unemp.youth    -0.601675  0.437154 -1.376  0.17798    
## unemp.adult     1.792263  0.856111  2.093  0.04407 *  
## median.assets  13.735847 10.583028  1.298  0.20332    
## num.low.salary  0.792933  0.235085  3.373  0.00191 ** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.94 on 33 degrees of freedom
## Multiple R-squared:  0.7692, Adjusted R-squared:  0.6783 
## F-statistic: 8.462 on 13 and 33 DF,  p-value: 0.0000003686

```

## Exploring the lm object

What kind of output do we get when we run a linear model (`lm`) in R?

```

# List all attributes of the linear model
attributes(crime.lm)

## $names
## [1] "coefficients"    "residuals"        "effects"          "rank"      
## [5] "fitted.values"   "assign"           "qr"              "df.residual"    
## [9] "contrasts"        "xlevels"          "call"             "terms"        
## [13] "model"            
## 
## $class
## [1] "lm"

```



```
# coefficients
crime.lm$coef
```

```
##          (Intercept)      young.males      is.south1      average.ed
## -691.837587905      1.039809653     -8.308312889    18.016010601
## exp.per.cap.1960  exp.per.cap.1959      labour.part male.per.fem
##      1.607818377     -0.667258285     -0.041031047    0.164794968
##      population      nonwhite      unemp.youth unemp.adult
##     -0.041276887      0.007174688     -0.601675298    1.792262901
## median.assets  num.low.salary
##      13.735847285      0.792932786
```

255



None of the attributes seem to give you p-values. Here's what you can do to get a table that allows you to extract p-values.

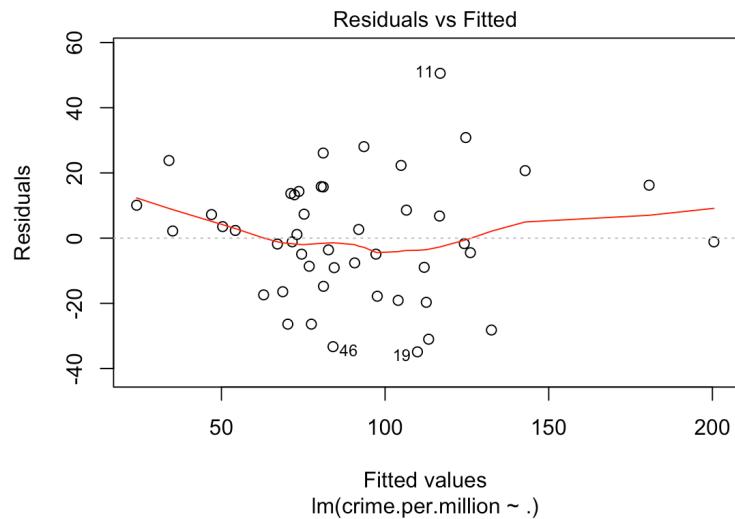
```
# Pull coefficients element from summary(lm) object
round(summary(crime.lm)$coef, 3)
```

	Estimate	Std. Error	t value	Pr(> t )
## (Intercept)	-691.838	155.888	-4.438	0.000
## young.males	1.040	0.423	2.460	0.019
## is.south1	-8.308	14.912	-0.557	0.581
## average.ed	18.016	6.497	2.773	0.009
## exp.per.cap.1960	1.608	1.059	1.519	0.138
## exp.per.cap.1959	-0.667	1.149	-0.581	0.565
## labour.part	-0.041	0.153	-0.267	0.791
## male.per.fem	0.165	0.210	0.785	0.438
## population	-0.041	0.130	-0.319	0.752
## nonwhite	0.007	0.064	0.112	0.911
## unemp.youth	-0.602	0.437	-1.376	0.178
## unemp.adult	1.792	0.856	2.093	0.044
## median.assets	13.736	10.583	1.298	0.203
## num.low.salary	0.793	0.235	3.373	0.002

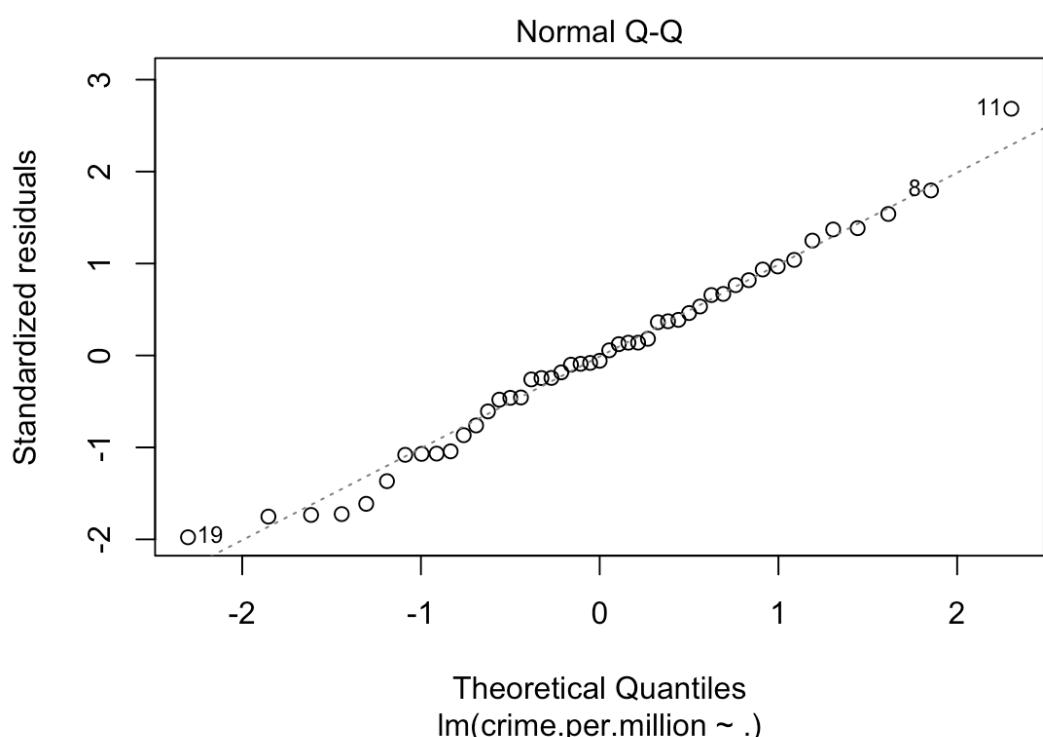
256

# Plotting the lm object

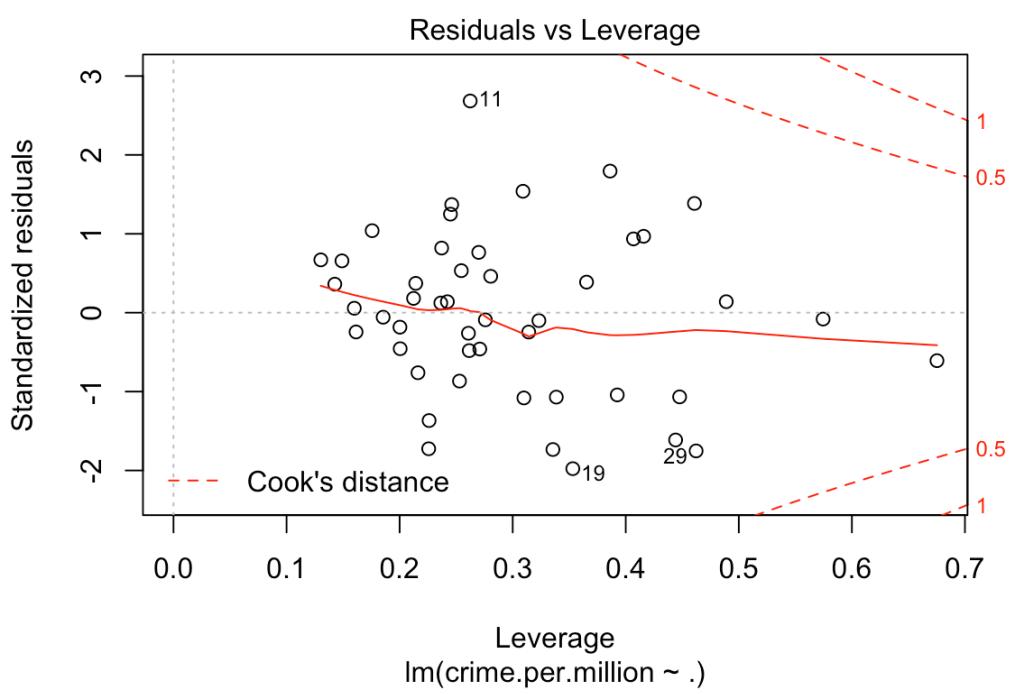
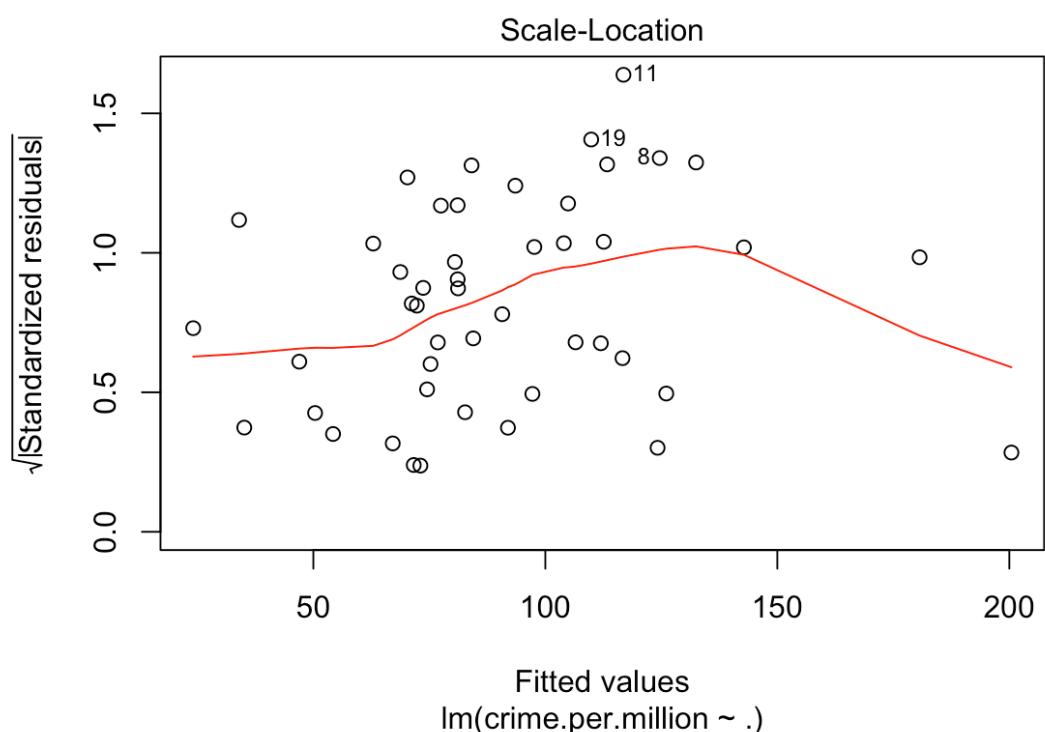
```
plot(crime.lm)
```



257



258



# Residuals vs Fitted Plot

Residuals vs. Fitted When a linear model is appropriate, we expect

- 1 the residuals will have constant variance when plotted against fitted values; and
- 2 the residuals and fitted values will be uncorrelated.

If there are clear trends in the residual plot, or the plot looks like a funnel, these are clear indicators that the given linear model is inappropriate.

261

# Normal QQ Plot

Normal QQ plot You can use a linear model for prediction even if the underlying normality assumptions don't hold. However, in order for the p-values to be believable, the residuals from the regression must look approximately normally distributed.

262

# Scale-Location Plot

Scale-location plot This is another version of the residuals vs fitted plot.  
There should be no discernible trends in this plot.

263

# Residuals vs Leverage

Residuals vs Leverage. Leverage is a measure of how much an observation influenced the model fit. It's a one-number summary of how different the model fit would be if the given observation was excluded, compared to the model fit where the observation is included. Points with high residual (poorly described by the model) and high leverage (high influence on model fit) are outliers. They're skewing the model fit away from the rest of the data, and don't really seem to fit with the rest of the data.

264

# Workshop 10: Regression

We'll begin by loading some packages.

```
library(MASS)  
library(plyr)
```

## Interaction terms in regression

```
# Building up the familiar birthwt data...  
  
# Rename the columns to have more descriptive names  
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",  
  "race", "mother.smokes", "previous.prem.labor", "hypertension", "uterine.irr",  
  "physician.visits", "birthwt.grams")  
  
# Transform variables to factors with descriptive levels  
birthwt <- transform(birthwt,  
  race = as.factor(mapvalues(race, c(1, 2, 3),  
    c("white", "black", "other"))),  
  mother.smokes = as.factor(mapvalues(mother.smokes,  
    c(0,1), c("no", "yes"))),  
  hypertension = as.factor(mapvalues(hypertension,  
    c(0,1), c("no", "yes"))),  
  uterine.irr = as.factor(mapvalues(uterine.irr,  
    c(0,1), c("no", "yes"))))  
)
```

265

# Workshop 10 : Regression

(a) Run a linear regression to better understand how birthweight varies with the mother's age and smoking status (do not include interaction terms).

(b) What is the coefficient of mother.age in your regression? How do you interpret this coefficient?

(c) How many coefficients are estimated for the mother's smoking status variable? How do you interpret these coefficients?

266

